

# An Empirical Study of GPU Kernel Performance Gaps in Modern Domain-Specific Languages

Anonymous Author(s)

## Abstract

Modern GPU domain-specific languages (DSLs) such as Triton and TileLang promise near-vendor-library performance at a fraction of the development effort, and are increasingly the lowering target for compilers and LLM-based kernel generators. Yet a developer who writes a *functionally correct* DSL kernel has no dependable way to know whether it is *performant*. We show that a kernel can pass the prevailing evaluation benchmarks and still run hundreds of times slower than the vendor-library baseline: across 22 kernels in five operator categories, naive-but-idiomatic DSL kernels that pass standard correctness gating range from competitive to more than two orders of magnitude slower than cuBLAS/cuDNN.

The benchmarks practitioners rely on do not surface this. Prevailing DSL and LLM-generated-kernel benchmarks (e.g., KernelBench and TritonBench) gate primarily on correctness—at most a weak speedup-over-eager check—and therefore admit performance-poor kernels, while under-covering DSLs, data types, shapes, and hardware and only partially preventing reward hacking. A benchmark comprehensive enough to close this gap is combinatorially infeasible.

As a pragmatic alternative, we (i) characterize how large and structured the hidden performance gap is and why it arises, separating user-space *authoring* causes (which we show are largely reducible) from *code-generation* and *library-maturity* causes (which are not); and (ii) distill the result into lightweight, partly baseline-independent evaluation heuristics and a small set of recurring optimization patterns that guide DSL kernel development when no ground-truth benchmark exists. We find that authoring idiom dominates the gap—a single reduction-primitive change recovers the normalization and reduction family to library parity or better (with one hardware-specific exception we diagnose)—while a residual structural gap persists on convolution, and that a roofline-anchored comparability check flags poor kernels that correctness-only benchmarks pass. Our kernel suite, profiling infrastructure, and analysis are publicly available at <https://anonymous.4open.science/r/GPUKernelPerformance-6132/>.

## Keywords

GPU kernels, domain-specific languages, Triton, TileLang, cuBLAS, cuDNN, empirical study, performance analysis, convolution

## ACM Reference Format:

Anonymous Author(s). 2026. An Empirical Study of GPU Kernel Performance Gaps in Modern Domain-Specific Languages. In . ACM, New York, NY, USA, 19 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

## 1 Introduction

Deep learning systems are now widely deployed, and their efficiency depends heavily on a small number of GPU kernels that dominate execution time [6, 21]. Frameworks have historically delegated these kernels to vendor libraries such as cuBLAS and cuDNN [3, 15], but modern architectures increasingly require fused operators, specialized kernels, and model-specific computations that fall outside the fixed functionality those libraries expose [25]. Achieving high performance therefore now depends on the ability to generate efficient *custom* kernels, and domain-specific languages (DSLs) such as Triton [16] and TileLang [26] have become the standard way to write them: they express kernels at the tile level while delegating memory movement, scheduling, and code generation to a compiler. Their reach extends well beyond hand-written code—Triton is the lowering target for `torch.compile` [1], and DSLs are increasingly the output of LLM-based kernel generators [17]—so the efficiency of DSL-generated kernels is now a first-order concern.

The difficulty is that a developer, or an automated generator, who produces a *functionally correct* DSL kernel has no reliable way to tell whether it is *performant*. A correct kernel may be competitive with the vendor library, or it may be orders of magnitude slower, and nothing in the source makes the difference visible. When a DSL-generated kernel is correct yet much slower than the corresponding library routine, the developer must determine whether the loss arises from the algorithm, the schedule, or the compiler—and current toolchains offer little diagnostic support for that distinction [12]. This makes performance deficits hard to predict and hard to correct, which is a software-engineering problem as much as a compiler one: a high-level kernel language is useful in practice only if its performance behavior is something developers can evaluate and reason about.

The benchmarks practitioners actually rely on do not close this gap. The prevailing benchmarks for DSL and LLM-generated kernels—KernelBench [17] and TritonBench [11]—gate on *correctness*, reporting performance as an outcome rather than enforcing it as an acceptance criterion. As a result they admit performance-poor kernels: in our measurements a

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

Conference'17, Washington, DC, USA

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-x-xxxx-xxxx-x/YYYY/MM

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

naive but idiomatic TileLang kernel passes the KernelBench-style correctness gate while running more than  $300\times$  slower than PyTorch, and the gate's one reward-hacking guard never fires because it only flags suspiciously *fast* kernels. These benchmarks also cover a narrow slice of the space—often a single DSL, data type, shape, and GPU per task—so a kernel that passes tells a developer little about whether it is well written. And the obvious remedy, a benchmark comprehensive enough to span every operator, data type, shape, and architecture, is combinatorially infeasible to build and maintain.

This paper studies that evaluation problem directly. We ask three questions: (RQ1) whether existing benchmarks can distinguish well-written DSL kernels from performance-poor ones, and where they fall short; (RQ2) for kernels that pass such benchmarks, how large and structured the hidden performance gap to vendor libraries actually is, and what causes it; and (RQ3) how to guide DSL kernel development when a comprehensive benchmark is infeasible. We evaluate Triton and TileLang against PyTorch with cuBLAS/cuDNN on 22 kernels spanning five operator classes—GEMM, attention, convolution, normalization, and element-wise/reduction—with hardware-counter profiling via Nsight Compute. Primary measurements span two datacenter architectures—the NVIDIA A100-SXM4-40GB (Ampere) and the NVIDIA GH200 (Grace Hopper)—with further cross-architecture replay on the A100-PCIE-40GB and H100-80GB-HBM3 reported in section 10.

Our findings are threefold. First, the evaluation gap is real and large: naive-but-correct kernels pass the gate while running hundreds of times slower than the library—the live shipped TileLang LayerNorm is  $323\times$  slower under locked clocks, with deliberately constructed worst-case kernels reaching  $1293\times$ —and the dominant existing benchmarks neither gate on performance nor cover enough of the space to catch it. Second, the hidden gap is highly structured. Much of it is an *authoring* artifact—a single reduction-primitive change (`T.serial`→`T.reduce`) with native-dtype I/O closes a  $1500\times$  authoring cliff (GH200; naive-to-optimized) on TileLang normalization and brings it to or above library parity—while a residual gap on convolution is structural, reflecting cuDNN's implicit-GEMM maturity (absent Winograd support contributes only  $\approx 2\text{--}3\%$ ). We separate causes by *where the fix belongs*: user-space authoring (RC0a), code generation (RC0b/RC1/RC2b), and library maturity. Third, absent a comprehensive benchmark, lightweight heuristics suffice to flag poor kernels: a comparability check against the library, anchored by a baseline-independent roofline fraction, cleanly separates well-written kernels (0.6–0.9 of roofline, at or above parity) from poor ones (below 0.2 of roofline, far from parity), and a small set of recurring optimization patterns guides the fixes.

## Contributions.

- **The evaluation gap (RQ1).** We show that prevailing DSL and LLM-generated-kernel benchmarks

admit performance-poor kernels—demonstrated end-to-end against the actual evaluator—and survey their coverage and reward-hacking gaps.

- **The hidden gap and its causes (RQ2).** We characterize the performance gap these benchmarks miss across 22 kernels and five operator classes on multiple GPU architectures, with hardware-counter root causes that separate authoring, code-generation, and library-maturity factors.
- **Guidance without a benchmark (RQ3).** We propose and validate lightweight evaluation heuristics—a library-comparability screen plus a baseline-independent roofline anchor—and a set of recurring optimization patterns that guide DSL kernel development and recover most of the gap.
- **Artifact.** We release the kernel suite, profiling infrastructure, and analysis to support reproducibility and follow-on work.

## 2 Background

### 2.1 Tile-Based GPU Programming

A GPU executes thousands of threads organized into *thread blocks*, each occupying a *streaming multiprocessor* (SM) with fast programmer-managed scratchpad (*shared memory*). The roofline model [27] captures the two performance limits: arithmetic throughput (TFLOPS) and memory bandwidth (GB/s); high-performance kernels tile data [28] to maximize shared memory reuse and stay compute-bound.

### 2.2 Triton

Triton [23] is a Python-embedded DSL and compiler in which the unit of computation is a *tile*: a statically shaped, multi-dimensional array operated on by all threads in a program instance. The Triton compiler performs memory coalescing, shared memory allocation, thread swizzling, and software pipelining automatically, targeting NVIDIA PTX through an MLIR-based compilation pipeline [8]. Triton exposes performance tuning via `@triton.autotune`, sweeping programmer-specified tile configurations and caching the fastest per problem shape. Triton has been integrated as the default lowering target for `torch.compile` since PyTorch 2.0 [1].

Triton's support for convolution is less mature: an `im2col`-style lowering exists but has been observed to fall substantially short of cuDNN performance on common workloads [20].

### 2.3 TileLang

TileLang [26] is a DSL developed at Microsoft Research that separates the *algorithm* (what to compute) from the *schedule* (how to map computation onto GPU resources) more explicitly than Triton, allowing the programmer to independently tune memory staging, thread layout, and pipeline depth without changing the functional description. TileLang compiles through Apache TVM's infrastructure with explicit hardware intrinsic support on NVIDIA and AMD architectures. TileLang's convolution support and its performance relative to

cuDNN have not been systematically evaluated prior to this work.

## 2.4 cuBLAS and cuDNN

cuBLAS [15] is NVIDIA’s BLAS implementation for CUDA, providing highly optimized routines for dense linear algebra, with cuBLASLt offering a flexible GEMM API with configurable layouts and compute types (FP32, FP16, BF16, FP8, INT8). cuBLAS dispatches the fastest GEMM implementation via a runtime heuristic recommender trained on problem shapes and hardware generation. The underlying template library, CUTLASS [13], is publicly available and serves as a reference for DSL performance comparisons.

cuDNN [3, 4] covers the full set of deep learning primitives: convolution (forward and backward passes), pooling, normalization, recurrent layers, and attention. cuDNN selects among implicit GEMM, Winograd ( $3 \times 3$  filters), FFT (large filters), and direct convolution at runtime via `cudaDnnFind`, benchmarking all candidates per configuration.

## 2.5 Benchmarking DSL and LLM-Generated Kernels

Custom GPU kernels are increasingly produced not by hand but by automated tools: `torch.compile` lowers operators to Triton [1], and a growing body of work uses LLMs to generate kernels directly [10, 17]. This has made kernel *benchmarks* the de-facto arbiters of kernel quality, and two dominate the DSL and LLM-generated setting.

TritonBench [11] is a collection of production-grade Triton kernels curated from high-starred open-source repositories, covering attention variants, GEMM, softmax, normalization, and fused operations. It reports both functional correctness and efficiency metrics—memory-bandwidth utilization and GPU efficiency as a fraction of theoretical peak TFLOPS—though for Triton only, on a single GPU, and for one data type and shape per operator. Our kernel suite is derived from TritonBench’s GitHub channel (184 kernels from 95 repositories), supplemented with TileLang re-implementations and additional convolution configurations.

KernelBench [17] is the de-facto benchmark for LLM kernel generation: each task poses a PyTorch reference module and accepts a generated kernel—CUDA in the original, and Triton/TileLang via the same harness in our artifact—that reproduces the reference output on randomized inputs. It records a speedup over the eager baseline but *gates on correctness alone*, and its sole reward-hacking guard flags only kernels that run implausibly *fast*. Neither benchmark rejects a kernel for being slow, and neither spans more than a narrow slice of the (DSL  $\times$  data type  $\times$  shape  $\times$  architecture) space—the evaluation gap we quantify in section 5.

## 3 Methodology

We evaluate whether current GPU DSLs can match vendor libraries on representative deep learning operators, and, when they do not, identify the cause of the gap. To this end, we construct a kernel suite that covers the main computation

patterns in modern models, compare DSL implementations against library-backed PyTorch baselines under matched configurations, and profile both end-to-end performance and low-level hardware behavior.

### 3.1 Kernel Suite

Our benchmark suite covers five operator categories that capture the dominant computation patterns in modern deep learning: matrix multiplication, attention, convolution, normalization, and element-wise or reduction operators. In total, the suite contains 22 kernels: three **GEMM** workloads (i.e., dense matrix multiplication, batched matrix multiplication, and fused linear plus activation), one **attention** workload (i.e., scaled dot-product attention / FlashAttention variants), one **convolution** workload (i.e., Conv2d with  $1 \times 1$  to  $7 \times 7$  filters, including depthwise and strided cases), two **normalization** workloads (i.e., LayerNorm and RMSNorm), and fifteen **element-wise or reduction** workloads including add, mul, relu, leaky\_relu, softmax, log\_softmax, logsumexp, swiglu, argmax, max reduction, mean reduction, cross-entropy, matrix transpose, index select, and embedding.

We derive these workloads from TritonBench [11], the TileLang example repository [26], and our own implementations for configurations not covered by either source. We exclude operators with sparse inputs or runtime-dependent output shapes, since they do not admit stable or meaningful library baselines.

*Selection criteria.* We narrow the kernels in TritonBench [11] to our 22-kernel suite using four criteria. First, *forward-pass only*: we exclude backward and gradient kernels, whose performance is governed by different access patterns and reduction structures. Second, *stable library baseline*: each kernel must admit a well-defined cuBLAS, cuDNN, or PyTorch reference that computes the same function, so that library efficiency is a meaningful quantity. Third, *five-category coverage*: we retain kernels that populate the five operator categories above rather than over-sampling any single class. Fourth, *canonical operators*: we exclude operators with sparse inputs or runtime-dependent output shapes, which lack a stable library baseline. The Triton implementations are drawn from TritonBench, except `layer_norm`, which follows the TorchInductor reference; the TileLang implementations are our re-implementations of the same operators, and any kernel absent from both source suites is implemented by us.

### 3.2 Baseline Construction

For each workload, we compare DSL implementations against the corresponding vendor-library path exposed through PyTorch. GEMM baselines use cuBLAS via `torch.matmul`. Convolution baselines use PyTorch `nn.Conv2d` with default NCHW memory layout under standard cuDNN algorithm selection. Normalization baselines use `F.layer_norm` and `F.rms_norm`. Element-wise and reduction baselines use standard PyTorch eager execution.

For attention, we use PyTorch scaled dot-product attention with the FlashAttention backend enabled through



`enable_flash_sdp(True)`. We report these results separately because the Triton and TileLang implementations are not algorithmically identical to the PyTorch baseline.

### 3.3 Profiling Setup

We measure both end-to-end latency and hardware performance counters. Latency is used for all comparisons against library baselines, while hardware counters are used to investigate the causes of observed performance gaps.

*End-to-end throughput.* For each workload and input configuration, we measure host-synchronized GPU execution time using `time.perf_counter()` bracketed by `torch.cuda.synchronize()` at entry and exit. Each measurement uses 10 warm-up iterations followed by 100 timed iterations, and we report the median over the timed runs. All results reflect GPU-side execution only and exclude host-side dispatch overhead. To reduce measurement noise, we run each workload in isolation without concurrent GPU activity. We pin the clocks for all timing at graphics 1215 MHz and memory 1215 MHz; the 1410 MHz maximum power-caps under the 400 W board limit, whereas 1215 MHz holds flat, yielding a run-to-run relative standard deviation of 0.0–0.9%.

*Hardware performance counters.* To support root-cause analysis, we collect hardware counter profiles with NVIDIA Nsight Compute [14]. The RQ2 analysis is grounded in seven concrete counters: `gld_efficiency` (global-load bytes-per-sector efficiency), `ldg_per_request` (sectors per load request), `n_regs` (registers per thread), `n_spills` (register spills), `warp_stall_long_scoreboard` (memory-latency stall cycles), `warp_stall_barrier` (synchronization stall cycles), and `lts_t_sector_hit_rate` (L2 sector hit rate). Each profile is collected from a single execution, and we verify stability within 2% across five repeated runs.

*Correctness.* Before any timing, we validate every DSL kernel against its library baseline using per-dtype numerical tolerances: `fp32` at `atol/rtol = 1e-5`, `fp16` at `1e-3`, and `bf16` at `1e-2`, with the tolerance relaxed by  $2\times$  for reduction kernels to absorb order-of-summation differences. We additionally exercise an edge-case suite (NaN, Inf, denormal, and all-equal inputs) across all 22 kernels, observing 0 crashes. After applying the RQ3 mitigations, we revalidate the fixed kernels and obtain 5/5 passes, confirming the optimizations preserve numerical equivalence. We exclude FP32 GEMM from the correctness-gated comparison because its discrepancy is root-caused as silent TF32 lowering on the SM80 tensor cores rather than a kernel defect (see section 6).

### 3.4 Metrics

Our primary metric is **library efficiency**,

$$E_{\text{lib}} = \frac{t_{\text{lib}}}{t_{\text{DSL}}} \times 100\%,$$

where  $t_{\text{lib}}$  denotes the latency of the cuBLAS or cuDNN baseline and  $t_{\text{DSL}}$  denotes the latency of the corresponding DSL implementation under the same hardware and input configuration. A value of 100% indicates parity with the library

baseline; lower values indicate that the DSL implementation is slower.

Secondary metrics: throughput ( $T = 2 \cdot \text{FLOPs}/t$ ), hardware efficiency ( $E_{\text{hw}} = T/T_{\text{peak}}$ ), and memory bandwidth utilization ( $B = \text{bytes}/(t \cdot B_{\text{peak}})$ )—are used qualitatively in root-cause analysis (RQ2) to interpret counter measurements; they are not tabulated separately.

### 3.5 Evaluation-Gap and Heuristic Measurement

To test whether existing benchmarks admit slow kernels (RQ1), we run each naive DSL kernel through a KernelBench-style evaluator (`eval_kernel_against_ref`), which overrides the candidate’s input generators with the reference’s, checks output equivalence on randomized inputs, and warns only when a kernel exceeds  $10\times$  the reference speed; we record the pass/fail verdict and the measured slowdown against the PyTorch baseline. To compute the roofline anchor (RQ3), we model each kernel’s essential work—bytes moved for memory-bound kernels (input plus output tensors) or FLOPs for compute-bound kernels—and divide the achieved work-rate by the relevant hardware peak (HBM bandwidth or FP16 Tensor-Core throughput) for the measured GPU; the peak assumed for each kernel is recorded alongside its measurement.

### 3.6 Experimental Setup

*Hardware.* We measure on two primary datacenter architectures. The first is an NVIDIA A100-SXM4-40GB GPU (Ampere, `sm_80`) with 108 SMs, 40 GB HBM2e memory ( $\sim 1.5$  TB/s peak memory bandwidth), a 40 MB L2 cache, and a 400 W power cap (NVIDIA Driver 610.43.02, CUDA 12.8), on which we measure the suite magnitude (section 5) and root-cause counters (section 6). The second is an NVIDIA GH200 (Grace Hopper, `sm_90`; 96 GB HBM3 at  $\sim 4.0$  TB/s,  $\sim 989.5$  TFLOP/s FP16 Tensor-Core, clocks locked at 1320/2619 MHz), on which we run the evaluation-gap demonstration (section 5.2) and validate the roofline heuristic (table 7). Per-architecture magnitudes are reported where they differ.

*Software.* Our software stack consists of PyTorch 2.8.0+cu128, Triton 3.4.0, TileLang 0.1.6.post1, bundled cuDNN, and Nsight Compute 2026.2.0. We pin all versions in a reproducibility manifest distributed with the benchmark suite. For all experiments, we run under default cuDNN flags and pre-warm the Triton auto-tuner before measurement.

## 4 Research Questions

We organize the study around three research questions. The first asks whether existing benchmarks can tell a well-written DSL kernel from a performance-poor one; the second characterizes and explains the performance gap those benchmarks fail to surface; the third asks how to guide kernel development when a comprehensive benchmark is infeasible.

**RQ1 (Evaluation gap).** Do existing DSL and LLM-generated-kernel benchmarks distinguish well-written kernels from performance-poor ones, and along which dimensions—DSL coverage, data type, input shape, hardware, and reward-hacking prevention—do they fall short?

**RQ2 (Hidden gap and its causes).** For kernels that pass existing correctness-style benchmarks, how large and structured is the performance gap to vendor libraries, and which authoring, code-generation, and library-maturity factors cause it?

**RQ3 (Guidance without a comprehensive benchmark).** Absent a comprehensive benchmark, can lightweight evaluation heuristics—comparability with the vendor baseline plus a roofline anchor—together with recurring optimization patterns reliably flag poor kernels and guide fixes? Is there a single dominant fix, and where are DSLs genuinely limited?

Our primary measurements span two datacenter architectures, the NVIDIA A100-SXM4-40GB (Ampere, [sm\\_80](#)) and the NVIDIA GH200 (Grace Hopper, [sm\\_90](#)); further cross-architecture replay on the A100-PCIE-40GB and H100-80GB-HBM3 is reported in section 10, confirming that the mechanisms identified in RQ2 transfer across GPU families.

## 5 The Evaluation Gap (RQ1)

This section answers RQ1: *do existing benchmarks distinguish well-written DSL kernels from performance-poor ones, and along which dimensions do they fall short?* We survey what the prevailing DSL and LLM-generated-kernel benchmarks measure (section 5.1), show that a correct-but-slow kernel passes their gate unflagged (section 5.2), and then quantify the hidden performance gap this gate admits across the 22-kernel suite (section 5.3 onward).

### 5.1 What Existing Benchmarks Measure

The two benchmarks that dominate DSL and LLM-generated GPU-kernel evaluation are KernelBench [17] and TritonBench [11], and both treat performance as a *reported outcome* rather than an *acceptance criterion*. KernelBench, the de-facto target for LLM kernel generation, accepts a candidate on *correctness alone*: its evaluator runs the kernel against a randomized-input PyTorch reference and the task passes if the outputs match within tolerance, with measured speedup recorded but never gated. Its only guard against reward hacking is a heuristic warning when a kernel runs more than 10× faster than the reference—so it fires only on suspiciously *fast* kernels and is silent on the slow ones, which are the failure mode we study here.<sup>1</sup> TritonBench is Triton-only and, to its credit, reports a roofline-anchored GPU-efficiency metric (achieved fraction of hardware peak)—a precedent we build on in section 7—but it likewise reports rather than gates on that metric, applies no anti-cheat, and fixes one data type

<sup>1</sup>We exercise this evaluator directly: the KernelBench-style harness in our artifact (`eval_kernel_against_ref`) overrides the candidate’s input generators with the reference’s and warns on more-than-10× speedups, the same anti-cheat logic shipped upstream.


and shape per operator on a single GPU. Neither benchmark covers more than a narrow slice of the (DSL × data type × shape × architecture) space, and neither rejects a kernel for being slow. A kernel that “passes” therefore certifies correctness, not quality.

### 5.2 Passing Is Not Fast

To show that this gap is not hypothetical, we run idiomatic but unoptimized DSL kernels—the kind a developer or an LLM generator would plausibly produce—through the KernelBench-style correctness gate and then time them. On the GH200, the live shipped TileLang LayerNorm kernel passes the gate on all five correctness shapes yet runs 323× slower than PyTorch at the large shape under locked clocks (51.5 ms vs. 0.16 ms); shipped TileLang RMSNorm passes and is 117× slower. A deliberately constructed worst-case LayerNorm variant—fp32 internal I/O and a zero-initialized output buffer—is slower still, reaching 1293× (176.4 ms, unlocked); even a naive argmax passes and is 16.7× slower. In every case the gate returns `correct=True` and the more-than-10× reward-hacking guard never triggers, because the kernels are slow, not fast. The shipped kernels are not adversarial constructions: they are the reduction kernels shipped as idiomatic TileLang, differing from a competent implementation only in the reduction primitive (section 6). The same authoring idiom shows the same qualitative gap on the A100 suite below, where the live shipped TileLang LayerNorm is 1087× slower (table 3; A100-SXM4, unlocked profile), with the suite magnitude and root-cause analysis measured on the A100 and the evaluation-gap and roofline results on the GH200; we report per-architecture magnitudes throughout and note where they differ.

### 5.3 The Hidden Gap These Benchmarks Admit

Figure 1 summarizes library efficiency (%) for all kernels in the suite, broken down by DSL (Triton, TileLang) and kernel category. The key finding is that the performance gap is not uniform: Triton is broadly competitive for element-wise and normalization kernels (88.7% of PyTorch throughput for LayerNorm and ~95–116% for element-wise kernels; the RMSNorm outlier at 851.8% is due to kernel fusion and is discussed separately in section 5.3) but trails substantially on convolution (19.3%) and large square GEMM (32.8%); TileLang achieves competitive throughput only on a narrow set of large shapes while exhibiting severe underperformance on reductions and normalizations (0.1% of PyTorch throughput on layer normalization). These figures are for the *idiomatic* kernels shipped in each DSL’s examples; we show in RQ3 (section 7) that the reduction and normalization gaps are authoring artifacts—a competent `T.reduce` rewrite recovers the family to PyTorch-comparable or better (with one hardware-specific exception we diagnose in section 7.5)—whereas the convolution and large-GEMM gaps largely persist even after best-effort optimization, marking the genuine residual.



**Figure 1: Library efficiency (%) of Triton and TileLang kernels relative to cuBLAS/cuDNN on the NVIDIA A100-SXM4-40GB, grouped by kernel category (log scale).** GEMM shows boxes (IQR; whiskers  $1.5 \times$  IQR); all other categories show individual kernel measurements as dots ( $n = 2$  per cell). Dashed line marks parity (100%). The hollow diamond at 851.8% (Triton, Norm.) marks Triton rms.norm: the PyTorch eager baseline does not fuse its normalization passes, making this an unfair comparison rather than a genuine DSL win; see section 5.3. The boxplot flier at 548.6% (TileLang, GEMM) marks fused linear+activation, a valid DSL advantage through kernel fusion; see section 5.3. Attention excluded; see section 5.3.

**Finding: On the A100-SXM4, Triton achieves 32.8–92.9% of PyTorch/cuBLAS throughput on GEMM depending on shape; TileLang achieves 7.0–80.0% with a complementary strength profile that reorders on the GH200 (table 4). TileLang achieves 548.6% on fused linear+activation by fusing two passes into one kernel, while Triton degrades to 32.8% on the 16384<sup>2</sup> square matmul where its auto-tune search space is under-populated.**

Table 1 reports latency (ms) and library efficiency for FP16 matrix multiplication across representative shapes.

*Notation:* 16384<sup>2</sup> denotes a square 16384  $\times$  16384 FP16 GEMM; 64  $\times$  128<sup>2</sup> denotes a batched GEMM of batch 64 over 128  $\times$  128 matrices. FP32 TileLang GEMM is excluded; the failure is diagnosed as silent TF32 lowering in `T.gemm` (the SM80 F32TF32F32F32 MMA atom truncates the mantissa to 10 bits; see section 6), not a logic error.

The 32.8% efficiency on the 16384<sup>2</sup> matmul reflects the auto-tuning search space limitation (RC2, section 6): Triton’s default configuration list for square GEMM does not cover optimal tile shapes at this scale, and register pressure from large tile configurations reduces SM occupancy. TileLang performs better on this shape (80.0%) because its explicit

**Table 1: GEMM-family kernels (FP16) latency (ms, lower is better), per configuration on the A100-SXM4 and GH200 (both unlocked tuned profile).  $E_{lib}$  = library efficiency relative to PyTorch/cuBLAS baseline. Bold marks best DSL result per row.**

| Configuration                                                      | PyTorch | Triton        | $E_{lib}$     | TileLang     | $E_{lib}$     |
|--------------------------------------------------------------------|---------|---------------|---------------|--------------|---------------|
| <i>Square matmul</i> 16384 <sup>2</sup>                            |         |               |               |              |               |
| A100                                                               | 34.24   | 104.34        | 32.8%         | <b>42.80</b> | <b>80.0%</b>  |
| GH200                                                              | 13.71   | 24.34         | 56.3%         | <b>22.96</b> | <b>59.7%</b>  |
| <i>Batched matmul</i> 64 $\times$ 128 <sup>2</sup>                 |         |               |               |              |               |
| A100                                                               | 0.0787  | <b>0.0959</b> | <b>82.1%</b>  | 0.1338       | 58.8%         |
| GH200                                                              | 0.0670  | <b>0.0566</b> | <b>118.4%</b> | 0.1059       | 63.3%         |
| <i>Batched matmul</i> 128 $\times$ 2048 <sup>2</sup>               |         |               |               |              |               |
| A100                                                               | 0.811   | <b>0.873</b>  | <b>92.9%</b>  | 11.55        | 7.0%          |
| GH200                                                              | 0.351   | 0.535         | 65.6%         | <b>0.318</b> | <b>110.4%</b> |
| <i>Fused linear+activation</i> 1 $\times$ 2048 $\rightarrow$ 16384 |         |               |               |              |               |
| A100                                                               | 36.03   | 49.71         | 72.5%         | <b>6.57</b>  | <b>548.6%</b> |
| GH200                                                              | 12.48   | 7.35          | <b>169.9%</b> | 7.66         | 163.0%        |

schedule specifies pipeline stages independently of tile size. For batched GEMM, both DSLs are competitive at the small scale (Triton 82.1%, TileLang 58.8%) but diverge at the large scale (Triton 92.9%, TileLang 7.0%), with TileLang further degraded by per-launch JIT compilation overhead. The fused linear+activation result shows TileLang achieving 548.6%—it fuses the two-pass computation into a single kernel, avoiding the memory-allocation overhead in PyTorch’s eager path; Triton (72.5%) does not achieve the same fusion benefit on this shape.

*Attention.* Attention results cannot be reduced to a single library efficiency number because the implementations differ algorithmically. Triton’s kernel implements FlashAttention-2 [5] (44.83 ms at the large configuration), which rewrites the attention computation to process the score matrix in tiles, avoiding materializing the full  $n \times n$  matrix; comparing this to PyTorch’s standard  $O(n^2)$  path (3045.0 ms) measures an algorithmic improvement, not DSL-versus-library productivity. TileLang’s kernel (5458.97 ms for the same configuration) uses the standard  $O(n^2)$  formulation without tiling, making it slower than even PyTorch’s unoptimized path because it does not exploit the memory hierarchy for attention. Neither result constitutes a valid DSL-versus-library efficiency measurement for the same computation; both are excluded from table 4.

**Finding: Triton-written Conv2d kernels achieve 9.8–28.7% of PyTorch/cuDNN throughput across standard 1  $\times$  1–7  $\times$  7 filters at the large shape (up to 46.3% at the small shape); TileLang Conv2d achieves 7.1–12.6%, a deficit compounded by both per-launch JIT compilation overhead and absent vectorization of strided spatial accesses. Both DSLs collapse to  $\leq 5.4\%$  on depthwise convolution, which they lower to  $N$  per-group GEMM launches.**



Table 2 reports library efficiency for representative Conv2d configurations across six filter shapes. Inputs follow the NCHW layout (PyTorch’s default arrangement).

At the large shape Triton efficiency on the standard  $1 \times 1$ – $7 \times 7$  filters falls monotonically from 28.7% ( $1 \times 1$ ) through 19.3% ( $3 \times 3$ ), 13.7% ( $5 \times 5$ ), to 9.8% ( $7 \times 7$ ); the gap widens precisely on the Winograd-ineligible larger filters (section 6). TileLang holds a narrow 7.1–12.6% band over the same filters. Depthwise convolution, which both DSLs lower to  $N$  per-group GEMM launches, collapses to  $\leq 5.4\%$  for Triton and to  $\sim 0.05\%$  for TileLang. Nsight Compute shows the conv gap is occupancy- and coalescing-bound (RC1 + RC2), not a register-spill problem: `n_spills = 0` at every filter size for both DSLs. On the GH200, Triton convolution stays comparably weak (12.1% at  $3 \times 3$ ,  $\leq 9.8\%$  at the larger filters), but TileLang convolution is far stronger than on the A100 (32.9–71.7% across  $1 \times 1$ – $7 \times 7$ ; table 2), so the TileLang convolution collapse is largely an Ampere-specific scheduling artifact rather than intrinsic to the DSL—the persistent residual is the *Triton* convolution gap, which holds on both architectures. The root causes of these gaps are analyzed in section 6.

**Finding: Triton normalization kernels match or substantially outperform PyTorch’s eager paths; TileLang normalization collapses to less than 1% of PyTorch throughput at the tested size ( $8192 \times 8192$ ), with LayerNorm  $1087\times$  slower.**

Table 3 reports latency for layer normalization and RMS normalization. Triton’s `layer_norm` achieves 88.7% of PyTorch throughput at the large shape ( $8192 \times 8192$ , BF16). Triton’s `rms_norm` achieves 851.8% efficiency because PyTorch’s eager `rms_norm` path does not fuse the variance-reduction and normalization passes, whereas the Triton kernel issues a single fused pass. TileLang normalization is catastrophically slow: 0.1% for LayerNorm (364.09 ms vs. PyTorch’s 0.335 ms, a  $1087\times$  slowdown) and 0.9% for RMSNorm (294.27 ms vs. 2.516 ms). As analyzed in section 6, this reflects a combination of per-launch JIT overhead, absent vectorized memory loads, and serialized memory-latency exposure under `T.serial` reduction loops (`warp_stall_long_scoreboard` dominates; `warp_stall_barrier`  $\approx 0$ ).

**Finding: Hardware-agnostic heuristic tuning shows negligible benefit for Triton in the tested configuration — default and tuned configurations produce nearly identical latencies for all kernels. The convolution gap is essentially unchanged (19.3%). TileLang shows no measurable benefit from tuning across any kernel.**

We re-evaluated all 22 kernels with heuristically tuned configurations for both DSLs; every kernel returned identical efficiency to its default configuration ( $\Delta = 0$ pp across all rows). This indicates that either the default configurations are already optimal for this hardware, or the tuner’s search space does not cover configurations that would improve performance on the NVIDIA A100-SXM4-40GB. The convolution gap (19.3%) is unchanged by tuning, consistent

with RC1–RC3 being structural compiler limitations rather than search-space coverage problems (section 6).

The apparent tension between this  $\Delta \approx 0$ pp result and the matmul speedup reported in section 7.4 is a search-space distinction, not a shape distinction. Here, “heuristic tuning” refers to the 12-config block-tile grid (varying only block- $M/N/K$ ); on the  $16384^2$  square matmul this grid yields  $\Delta \approx 0$ pp (Triton 32.5% untuned vs. 32.5% tuned). The expanded search in section 7, which additionally tunes `GROUP_SIZE_M`, `num_warps`, and `num_stages`, recovers Triton from 32.5% to 81.8% at  $16384^2$  (a  $2.52\times$  latency improvement) and from 56.4% to 76.5% at  $4096^2$  ( $1.36\times$ ); the optimal configuration is simply absent from the default block-tile grid (RC2a).

Table 4 summarizes median library efficiency per category and DSL.

The most striking asymmetry is TileLang’s normalization collapse: a  $1087\times$  slowdown on large LayerNorm (section 6) attributable to three compounding deficiencies (RC0), and it persists across architectures ( $329\times$  on the GH200, table 3). On the A100, Triton is broadly competitive with PyTorch for element-wise ( $\sim 95$ – $116\%$ ) and normalization (88.7%) kernels but exhibits a persistent gap on convolution (19.3%) and large square GEMM (32.8%). TileLang on the A100 shows competitive throughput only on fused linear+activation (548.6%) and large square matmul (80.0%); all other A100 TileLang results trail PyTorch substantially, with element-wise operations at  $\sim 67$ – $74\%$  and batched GEMM at 7.0%. The GH200 reorders this strength profile (table 4): TileLang is markedly more competitive there—batched GEMM 110%, convolution 54%, square matmul 60%, overall median  $\sim 67\%$  versus  $\sim 29\%$  on the A100—so the two gaps that survive on *both* architectures are the TileLang normalization collapse and the Triton convolution residual ( $\leq 13\%$  at  $3 \times 3$ ), not a fixed per-DSL ranking. This hardware-dependence is itself the point: a DSL kernel’s competitiveness cannot be read off one GPU.

Table 5 reports per-kernel library efficiency for the 15 element-wise and reduction kernels individually at the large input size. Entries above 100% (marked  $^\dagger$ ) reflect kernel fusion or an unfused/eager PyTorch baseline rather than genuine speedups over a vendor library, and should not be read as pure DSL wins.

## 6 Root Causes of the Hidden Gap (RQ2)

The kernels in section 5 pass existing correctness benchmarks (section 5.1) yet lag the vendor library by the margins reported there. This section identifies the causes of that hidden gap, separated by where the fix belongs: user-space *authoring* (RC0a), compiler *code generation* (RC0b, RC1, RC2a), and *library maturity*—the tuning and algorithm selection a mature vendor library provides (RC2b, RC4). This separation is what makes the gap actionable: authoring causes a developer can fix today, code-generation causes require compiler work, and library-maturity causes bound what any DSL kernel can currently reach.

*The following root causes are motivated by latency measurements (section 5) and community reports [20].*

**Table 2: Conv2d library efficiency ( $E_{\text{lib}}$ , %) on the A100-SXM4 (small and large shapes) and GH200 (large shape; unlocked tuned profile). NCHW layout (PyTorch default), FP16, padding “same”.  $E_{\text{lib}}$  relative to PyTorch/cuDNN baseline (higher is better). Bold marks the best DSL per cell.**

| Filter                 | A100 $8 \times 64 \times 56^2$ |          | A100 $32 \times 256 \times 128^2$ |          | GH200 $32 \times 256 \times 128^2$ |               |
|------------------------|--------------------------------|----------|-----------------------------------|----------|------------------------------------|---------------|
|                        | Triton                         | TileLang | Triton                            | TileLang | Triton                             | TileLang      |
| $1 \times 1$           | <b>44.09%</b>                  | 5.28%    | <b>28.69%</b>                     | 7.15%    | 28.97%                             | <b>32.89%</b> |
| $3 \times 3$           | <b>44.87%</b>                  | 4.12%    | <b>19.33%</b>                     | 12.62%   | 12.14%                             | <b>54.05%</b> |
| $5 \times 5$           | <b>30.19%</b>                  | 7.58%    | <b>13.67%</b>                     | 11.16%   | 8.08%                              | <b>49.48%</b> |
| $7 \times 7$           | <b>20.53%</b>                  | 9.05%    | <b>9.79%</b>                      | — (OOM)  | 9.77%                              | <b>71.70%</b> |
| depthwise $3 \times 3$ | <b>5.40%</b>                   | 0.05%    | <b>1.16%</b>                      | 0.05%    | <b>0.84%</b>                       | 0.04%         |
| $3 \times 3$ stride 2  | <b>46.34%</b>                  | 4.17%    | <b>17.17%</b>                     | 16.77%   | <b>13.28%</b>                      | 7.24%         |

PyTorch/cuDNN A100 large-shape ( $32 \times 256 \times 128 \times 128$ ) baseline latencies (ms):  $1 \times 1$  0.518,  $3 \times 3$  3.445,  $5 \times 5$  7.716,  $7 \times 7$  13.846, depthwise 0.776, stride-2 1.261. TileLang  $7 \times 7$  OOMs on the A100 (40 GB) but not on the GH200, where it reaches 71.7%.

**Table 3: Normalization latency (lower is better).  $E_{\text{lib}}$  relative to PyTorch eager baseline.**

| Kernel    | Shape                         | PyTorch(ms) | Triton(ms)   | $E_{\text{lib}}$ | TileLang(ms) | $E_{\text{lib}}$ |
|-----------|-------------------------------|-------------|--------------|------------------|--------------|------------------|
| LayerNorm | 8192 <sup>2</sup> BF16, A100  | 0.335       | <b>0.378</b> | <b>88.7%</b>     | 364.09       | 0.1%             |
| LayerNorm | 8192 <sup>2</sup> BF16, GH200 | 0.140       | <b>0.188</b> | <b>74.8%</b>     | 46.13        | 0.3%             |
| RMSNorm   | 8192 <sup>2</sup> FP16, A100  | 2.516       | <b>0.295</b> | <b>851.8%</b>    | 294.27       | 0.9%             |
| RMSNorm   | 8192 <sup>2</sup> FP16, GH200 | 1.276       | <b>0.118</b> | <b>1081%</b>     | 38.87        | 3.3%             |

A100 rows are the unlocked tuned profile; locked-clock re-timing (graphics 1215 MHz / memory 1215 MHz) over 100 runs shows run-to-run relative std-dev 0.0–0.9% on near-parity kernels, so the gaps are statistically significant. GH200 rows are the unlocked tuned profile; TileLang LayerNorm there is 0.3% (329 $\times$ ), matching the locked-clock 323 $\times$  in section 5.2 within run-to-run noise.

*RC0: TileLang Compiler-Level Deficiencies.* Two mechanisms contribute to RC0, which we separate by where the fix lives: RC0a is a user-space kernel-authoring choice, while RC0b is a compiler-codegen deficiency.

*RC0a: TileLang Reduction Authoring ( $T.\text{serial} \rightarrow T.\text{reduce}$ ).*

First, the original TileLang normalization kernels implement the reduction over the normalized dimension using `T.serial` loops—TileLang’s sequential for-loop construct—iterating over each partial sum one at a time. For LayerNorm, which accumulates 8 bfloat16 partial statistics per thread (sum and sum-of-squares for the mean and variance passes), this means the reduction is computed in 8 sequential steps with no inter-thread parallelism. The native alternative, `T.reduce`, lowers to a parallel butterfly reduction via `tl::AllReduce`. A *butterfly reduction* is a logarithmic communication pattern in which threads exchange partial results in  $\log_2 N$  rounds (here  $\log_2(256/32)$  rounds over 256 threads grouped into warps of 32). This algorithmic structure, however, is not where the cost lies: on the A100, the performance gain from `T.reduce` comes from better memory-latency hiding—`warp_stall.long_scoreboard` drops from 59.56 cycles to  $\approx 0$ —not from removing barriers, since `warp_stall.barrier` is already  $\approx 0$  in both the `T.serial` and `T.reduce` kernels. The GH200 ncu counters

confirm the same signature on `sm_90: warp_stall.barrier=0` with `warp_stall.long_scoreboard=49.95` dominating, alongside a 34 GB local-store register spill (RC3). Replacing the `T.serial` loop with a single `T.reduce` call thus removes the dominant latency source: on the A100, LayerNorm drops from 364 ms to 0.270 ms (1347 $\times$  improvement, section 7.2).

*RC0b: Absent Vectorized Loads (Compiler Codegen).* Second, TileLang’s compiled normalization kernels do not issue 128-bit vectorized loads (`LDG.128`), instead fetching 16-bit bfloat16 elements individually. For a tensor of shape  $8192 \times 8192$ , this produces an excessive number of discrete memory transactions that bottleneck the memory controller and prevents the L2 cache from streaming coalesced data to the SMs. On the A100, NCU confirms this scalar-granularity access for TileLang LayerNorm: a load efficiency of 50% bytes/sector at one sector per request, leaving DRAM throughput at 59.2%. Unlike RC0a, this deficiency cannot be addressed in user space and requires a compiler-codegen fix.

*TileLang float32 lowering to TF32.* A third deficiency surfaced as a correctness failure under float32 precision but is rooted in silent TF32 lowering rather than a logic error. On the A100, TileLang’s `T.gemm` kernel fails an exact float32 check (atol  $10^{-5}$ ) on 31.8% of output elements, with



**Table 4: Median library efficiency (%) per kernel category and DSL on the A100-SXM4 and GH200 (both unlocked tuned profile), computed over large-size benchmark configurations. Attention excluded from Overall; see text.**

| Kernel category              | A100-SXM4                                  |                         | GH200       |             |
|------------------------------|--------------------------------------------|-------------------------|-------------|-------------|
|                              | Triton                                     | TileLang                | Triton      | TileLang    |
| GEMM                         | 32.8–92.9%                                 | 7.0–80.0%               | 56–170%     | 60–163%     |
| Attention                    | <i>see text — baselines non-equivalent</i> |                         |             |             |
| Convolution                  | 19.3%                                      | 12.6%                   | 12.1%       | 54.1%       |
| Normalization                | 88.7% <sup>‡</sup>                         | 0.1–0.9% <sup>‡</sup>   | 74.8%       | 0.8–9.9%    |
| Element-wise                 | ~95–116%                                   | ~67–74%                 | ~66–98%     | ~62–68%     |
| <b>Overall (excl. attn.)</b> | <b>~73%<sup>§</sup></b>                    | <b>~29%<sup>§</sup></b> | <b>~98%</b> | <b>~67%</b> |

<sup>‡</sup> TileLang LayerNorm (large shape) is 0.1% (1087× slower than PyTorch) due to serialized memory-latency exposure under `T.serial` (`warp_stall_long_scoreboard` dominates) and absent vectorized loads; see section 6. RMSNorm Triton efficiency is 851.8% because the PyTorch eager path does not fuse normalization passes.

<sup>§</sup> Overall figures are medians excluding attention; the Triton geomean (~126%) is inflated by fusion and eager-baseline outliers (`rms_norm`, `leaky_relu`, `cross_entropy`, `linear+act`), so we report the more representative median. Splitting the baseline into vendor cuBLAS/cuDNN, `torch.compile(max-autotune)` fused, and eager PyTorch (`fused_baselines`) confirms the convolution gap persists after fusion.

**Table 5: Per-kernel library efficiency ( $E_{lib}$ , %) for the 15 element-wise and reduction kernels at the large input size, NVIDIA A100-SXM4-40GB. <sup>†</sup> marks entries >100%, which reflect kernel fusion or an unfused/eager PyTorch baseline (not pure DSL speedups over a vendor library).**

| Kernel           | Triton               | TileLang             |
|------------------|----------------------|----------------------|
| add              | 89.5%                | 74.5%                |
| mul              | 69.7%                | 67.6%                |
| relu             | 95.4%                | 68.3%                |
| leaky_relu       | 427.9% <sup>†</sup>  | 1241.7% <sup>†</sup> |
| swiglu           | 267.1% <sup>†</sup>  | 104.2% <sup>†</sup>  |
| softmax          | 116.0% <sup>†</sup>  | 19.0%                |
| log_softmax      | 44.9%                | 17.0%                |
| logsumexp        | 475.7% <sup>†</sup>  | 57.6%                |
| argmax           | 4.7%                 | 5.2%                 |
| max_reduction    | 12.7%                | 4.7%                 |
| mean_reduction   | 72.5%                | 7.2%                 |
| cross_entropy    | 1132.2% <sup>†</sup> | 111.0% <sup>†</sup>  |
| embedding        | 441.9% <sup>†</sup>  | 70.1%                |
| index_select     | 60.5%                | 55.2%                |
| matrix_transpose | 340.1% <sup>†</sup>  | 107.4% <sup>†</sup>  |

a maximum absolute difference of 0.185 at the tested shape  $4096 \times 2048 \times 1024$ . This error is in the same magnitude class as cuBLAS’s own TF32 path, and a manual FMA

float32 kernel passes the same  $10^{-5}$  check—isolating the fault to `T.gemm`. The mechanism is that TileLang exposes no user-space TF32 knob: float32 `T.gemm` lowers to the SM80 F32TF32TF32F32 MMA atom, whose 10-bit input mantissa (vs. FP32’s 23-bit) truncates precision. The behavior is reproducible across four isolation arms in `exp_fp32_gemm.py`. For this reason, all TileLang GEMM measurements in this paper use FP16; float32 TileLang GEMM results are excluded. This is a precision-lowering artifact, not a performance gap, but has implications for training workloads that require true FP32 accumulation precision.

**Finding RC0: TileLang’s compiled kernels exhibit two structural performance deficiencies: (RC0a) `T.serial` reductions serialize memory-latency hiding rather than synchronization (`warp_stall_long_scoreboard` dominates; `warp_stall_barrier`  $\approx 0$ ), correctable in user space via `T.reduce`; and (RC0b) absent 128-bit vectorized memory loads for normalization, which require a compiler-codegen fix. These compound the per-kernel performance gaps measured in RQ1. A third deficiency—silent TF32 lowering of `T.gemm` under FP32—is correctness-affecting but orthogonal to performance.**

*RC1: Absent Vectorization of Strided Memory Accesses.* NVIDIA Ampere’s memory subsystem achieves peak bandwidth only when global memory loads are issued as 128-bit vector instructions (`LDG.128`). On Hopper, the Tensor Memory Accelerator (TMA) similarly requires aligned, contiguous access descriptors. For GEMM kernels, Triton’s compiler successfully emits vectorized loads because the data layout is row-major and tile boundaries are aligned to 16-byte boundaries.

Convolution complicates this in two ways. First, each output element requires gathering input values from a ( $K_H \times K_W$ )-window strided by (`stride_h`, `stride_w`) in the spatial dimensions. The resulting access pattern is a strided gather across the innermost dimension when the input is stored in NCHW layout (PyTorch’s default, as benchmarked), breaking the alignment requirements for `LDG.128`. Second, for filter sizes  $K_H \times K_W > 1$ , the compiler must prove at compile time that successive iterations of the spatial reduction loop produce aligned addresses; this proof is not discharged by Triton’s current alias analysis, so it conservatively emits scalar 32-bit loads.

The consequence is a cascade noted in the community [20]: un-vectorized accesses prevent `cp.async` from firing, because CUDA’s asynchronous copy instructions require the source address to produce a 128-bit aligned transfer. Without `cp.async`, the software pipeline cannot overlap data movement with computation, and the kernel stalls on global memory latency at every tile boundary. On the A100’s HBM2e interface ( $\approx 1.5$  TB/s peak), scalar 32-bit loads achieve only a small fraction of peak bandwidth; `LDG.128` vectorized loads are required to saturate the memory bus.

**Finding RC1: Triton-generated convolution code fails to issue vectorized (128-bit) load instructions,**

preventing asynchronous copy (`cp.async`) from being applied and collapsing the software pipeline.

NCU confirms the absent vectorization on the A100: Triton conv2d achieves a load efficiency of only 12.55% bytes/sector at 16.4 sectors per request, versus the `cp.async`-staged matmul path.<sup>2</sup>

*RC2a: Configurations Absent From the Default Grid.* Triton’s auto-tuner sweeps a programmer-specified list of configurations

{(BLOCK\_M, BLOCK\_N, BLOCK\_K, num\_stages, num\_warps)}.

For GEMM, this 5-dimensional space is well-understood: large square tiles ( $128 \times 128$ ) with 4–8 pipeline stages are optimal across most A100 problem shapes [23]. The reference Triton GEMM tutorial ships a configuration list that achieves near-cuBLAS performance for common shapes, but omits the L2-swizzle and large-shape configurations needed at scale.

This is evidenced by the large-matmul result in section 5.3: Triton achieves only 32.8% of cuBLAS throughput on a  $16384^2$  FP16 matrix multiplication (104.34 ms vs. PyTorch’s 34.24 ms), a shape not covered by standard tutorial configuration lists, compared to 92.9% on a  $128 \times 2048^2$  batched GEMM where the auto-tuner has well-populated configurations. The convolution gap (19.3% for the  $32 \times 256 \times 128 \times 128$ ,  $3 \times 3$  configuration) follows the same pattern at an even larger scale. Expanding the search beyond the default 12-config block-tile grid (sweeping `GROUP_SIZE_M`, `num_warps`, and `num_stages`) recovers Triton from 32.5% to 81.8% of cuBLAS at  $16384^2$  (a  $2.52 \times$  gain) and from 56.4% to 76.5% at  $4096^2$ , confirming that the missing configurations—not an exhausted search budget—bound default performance (section 7).

*RC2b: L2 Cache Residency.* For the  $16384^2$  matmul, the combined A, B, C footprint ( $\approx 1.6$  GB) vastly exceeds the A100’s 40 MB L2 cache; cuBLAS employs hardware-specific cache-blocking (via cuBLASLt’s plan selector) while Triton’s generic `tl.dot` compilation lacks equivalent heuristics, producing the  $\approx 3 \times$  latency penalty observed (34.24 ms vs. 104.34 ms). On the A100, NCU confirms the residency divide: Triton’s  $16384^2$  matmul reaches only a 49.4% L2 hit rate and is DRAM-bound at 85.9% DRAM throughput, whereas cuBLAS holds an 80.5% L2 hit rate at 28.8% DRAM throughput.

Convolution adds filter-size degrees of freedom:  $1 \times 1$  convolutions behave like GEMM (large tiles preferred), while  $3 \times 3$  require smaller  $K$ -dimension tiles to keep shared memory within 48 KB per SM. Default configuration lists do not cover  $5 \times 5$  and  $7 \times 7$  filters, leaving substantial performance on the table. TileLang additionally requires `num_stages` to be declared ahead of time; for large filters the optimal value depends on register pressure from spatial accumulation—a dependency current scheduling heuristics do not model.

<sup>2</sup>Triton matmul and cuBLAS read 0% on this load counter because their `cp.async/LDGSTS` staging bypasses the LDG instruction the counter tracks.

**Finding RC2: The static tile-configuration search space in Triton’s `@triton.autotune` omits configurations needed at scale (RC2a: L2-swizzle and large-shape tiles absent from the default grid), and Triton’s generic compilation lacks the L2 cache-blocking heuristics cuBLAS applies once the working set exceeds L2 (RC2b).**

*RC3: TileLang LayerNorm Register Spill (Not a Convolution Problem).* Each streaming multiprocessor on the A100 exposes a 65,536 32-bit register file (a budget Ada’s sm.89 shares), so a block of 128 threads at 64 registers per thread occupies it fully, preventing a second resident block and eliminating pipeline interleaving. Contrary to the submitted-PDF framing, NCU shows convolution does not hit this wall: Triton conv2d reports `n_spills=0` at every tested filter ( $1 \times 1$  through  $7 \times 7$ ), even though register usage rises with  $K^2$  (up to 224 regs/thread at large shapes). The spill is instead TileLang-LayerNorm-specific: the large LayerNorm kernel uses 254 registers per thread, drops to 12.46% achieved occupancy, and spills 206 GB of local-load plus 137 GB of local-store traffic. This register pressure—not any convolution effect—is what collapses occupancy below the level required for latency hiding.

**Finding RC3: The TileLang LayerNorm kernel exceeds the per-thread register budget (254 regs/thread), spilling to local memory (206 GB local-load + 137 GB local-store) and reducing achieved occupancy to 12.46%; convolution kernels, by contrast, do not spill (`n_spills=0` at all filters).**

*RC4: Absence of Algorithm-Level Diversity.* cuDNN’s runtime selector chooses Winograd for  $3 \times 3$  filters when arithmetic reduction is the bottleneck. Winograd convolution reduces the multiply-accumulate count by a factor of approximately  $2.25 \times$  for  $3 \times 3$  (Winograd  $F(2, 3)$ ) filters, at the cost of additional data transformation passes. Neither Triton nor TileLang expose the Winograd transformation as a schedulable primitive, so they cannot exploit this arithmetic reduction. Isolating this effect on the A100, however, bounds its size: a deterministic A/B comparison of cuDNN on  $3 \times 3$  stride-1 convolution differs by only 0.06% (a ratio of 0.9994), so absent Winograd selection accounts for at most  $\approx 2$ –3% of the gap. Moreover, the measured Triton-vs-cuDNN gap widens on Winograd-ineligible filters ( $3 \times 3$  s1 19.25%,  $3 \times 3$  s2 16.94%,  $5 \times 5$  13.5%,  $7 \times 7$  9.81%), indicating that the residual is general cuDNN implicit-GEMM tuning rather than algorithm selection.

**Finding RC4: Absent Winograd selection accounts for at most  $\approx 2$ –3% of the residual  $3 \times 3$  convolution gap on the A100; the remainder reflects cuDNN’s general implicit-GEMM tuning, and the gap widens on Winograd-ineligible filters.**

## 6.1 Summary of Root Causes

Table 6 maps each root cause to the kernel category it primarily affects, its estimated contribution to the throughput

gap, and the Nsight Compute counter most diagnostic for its detection.

## 7 Guidance Without a Comprehensive Benchmark (RQ3)

This section answers RQ3: *absent a comprehensive benchmark, can lightweight evaluation heuristics and a small set of recurring optimization patterns reliably flag performance-poor DSL kernels and guide their repair?* RQ1 (section 5) showed that a comprehensive performance benchmark is infeasible and that correctness gates admit slow kernels; RQ2 (section 6) explained why the gaps arise. We now distill that evidence into developer-facing guidance: a two-part *evaluation heuristic* for deciding whether a kernel is well written (section 7.1), and the *optimization patterns* that repair the gaps it flags (section 7.2 onward), validated by optimizing kernels across the suite on the A100-SXM4 and GH200. The central result is a clean dichotomy: most of the dramatic gaps are *authoring artifacts* that one dominant pattern fully repairs, while a minority are *genuine residuals* that survive best-effort optimization—and the heuristic tells the two apart without a ground-truth benchmark. We position the user-space fixes (RC0a, RC2a, the matmul L2-swizzle) as lower bounds on automatable recovery; RC0b and the broader RC2b search space require compiler support.

### 7.1 Evaluation Heuristics: Is This Kernel Well Written?

We propose two complementary checks a developer can apply without a curated benchmark.

*A comparability screen (pragmatic, baseline-relative).* A well-written DSL kernel should (i) come within a small factor of the vendor/PyTorch baseline on representative shapes, (ii) be at least at parity—ideally faster—at large input sizes where launch and framework overheads amortize, and (iii) show no catastrophic per-shape collapse. This screen is cheap and catches gross failures: the idiomatic TileLang LayerNorm that runs more than 300× slower (section 5.2) fails it immediately. Its limitation is circularity—PyTorch is simultaneously the baseline and the de-facto definition of “good,” so the screen cannot certify a kernel that merely matches an already-slow baseline, and it can be fooled into crediting a kernel that beats an *unfused* eager path for the wrong reason (the RMSNorm 851% “win” in table 3 is a fusion artifact, not headroom). It is therefore a screen, not a certificate.

*A roofline anchor (baseline-independent).* To break that circularity we anchor quality to the hardware rather than to PyTorch: for a kernel whose essential work is  $W$  (bytes moved if memory-bound, FLOPs if compute-bound) and whose optimized time is  $t$ , the achieved roofline fraction is  $\rho = W/(t \cdot P)$ , where  $P$  is the relevant hardware peak [27]. TritonBench reports an analogous achieved-peak metric [11]; we use it not as a leaderboard score but as a *decision signal*—a kernel near the memory or compute roof is well written regardless of what the library does. Table 7 reports  $\rho$  for

the optimized kernels on the A100-SXM4 (primary) and the GH200, alongside their PyTorch-relative efficiency  $E_{\text{lib}}$  and the *cliff* (naive-to-optimized speedup) that measures how much authoring headroom each kernel hid; the two architectures agree on every authoring-artifact vs. structural-residual classification (the GH200  $\rho$  values cited below are matched within  $\pm 0.1$  on the A100). The anchor does two things the comparability screen cannot. First, it confirms the genuine residuals are genuinely far from the hardware: optimized Triton Conv2d reaches only  $\rho = 0.16$  and index-returning argmax only  $\rho = 0.10$ , so their shortfall is real headroom, not merely a fast baseline. Second, it de-inflates baseline-relative outliers: the RMSNorm kernel that looks like a 13× “win” over PyTorch sits at  $\rho = 0.69$ —well written, but not 13× beyond what the hardware allows. The recovered normalization and reduction family lands at  $\rho = 0.41$ – $0.87$ , while the residual convolution and argmax kernels sit at  $\rho \leq 0.28$  even after best-effort optimization; the cliff column shows these are not the same axis—LayerNorm hid a 1541× authoring artifact that the dominant pattern fully recovers, whereas Conv2d’s 8.2× cliff still leaves it at  $\rho = 0.28$ , a structural ceiling the rewrite cannot lift. We present the two checks together because neither suffices alone, and they disagree exactly in a judgment band near  $\rho \approx 0.5$ : Softmax and Matmul both achieve  $\rho = 0.47$ , yet Softmax is at parity ( $E_{\text{lib}} = 0.87$ ) while Matmul trails cuBLAS ( $E_{\text{lib}} = 0.68$ )—only the two checks together make the call.

### 7.2 Normalization Kernels: Correcting RC0

*Finding: Correcting RC0 with two optimizations (each necessary, neither sufficient alone) brings TileLang LayerNorm to 124% and RMSNorm to 990% of PyTorch throughput on the A100-SXM4— a 1347× and 1157× reduction over the unoptimized TileLang baseline (table 3): (1) replacing *T.serial* reduction loops with the native *T.reduce* primitive (the dominant step; in the RTX 4000 Ada optimization campaign this change alone gave 208×); (2) switching to native bfloat16/float16 I/O and eliminating intermediate .float() precision casts. These results confirm RC0 as the primary cause of TileLang’s normalization deficit and as fully correctable in user space.* Here one iteration denotes a single kernel-source edit plus one benchmark run with a correctness check, logged and committed in ITERATIONS.md; failed and regressing attempts count. The LayerNorm ITERATIONS.md records 18 such iterations and the RMSNorm log records one.

*Step 1: Replace T.serial with T.reduce:* Replacing the manual *T.serial* loop with a single *T.reduce* call is the dominant fix. Subsequent iterations explored thread count, tiling, and two-pass formulations, confirming the loop structure— not thread configuration—was the bottleneck.

*Step 2: Native bfloat16 I/O:* Switching to native bfloat16 tensors and removing intermediate .float() casts brings the kernel close to the memory-bandwidth limit: the theoretical minimum for a  $8192 \times 8192$  bfloat16 tensor (read + write  $\approx 268$  MB) at the A100’s  $\approx 1.5$  TB/s is  $\approx 0.18$  ms; the



**Table 6: Root-cause summary.** Measured contribution is the latency speedup recovered by the corresponding mitigation (section 7); “—” indicates no direct mitigation experiment conducted. RC0a, RC0b, and RC3 are TileLang-specific; RC1, RC2, and RC4 affect both DSLs. Per-RC counter values are consolidated in the A100 and GH200 `ncu_summary.csv` files, which agree on the RC signatures.

| Root cause                                                   | Affected kernels              | Contribution             | Diagnostic counter                                                            |
|--------------------------------------------------------------|-------------------------------|--------------------------|-------------------------------------------------------------------------------|
| RC0a: <code>T.serial</code> instead of <code>T.reduce</code> | All TileLang reductions, norm | 1347× (Layer-Norm, A100) | <code>warp_stall_long_scoreboard</code> (barrier $\approx 0$ )                |
| RC0b: Absent vectorized loads / dtype cast                   | TileLang norm                 | —                        | <code>l1tex</code> bytes/sector                                               |
| RC1: Strided access, scalar LD                               | Conv2d ( $K > 1$ ), Triton    | 2.57× with RC2           | Load bytes/sector eff. (12.55%)                                               |
| RC2a: Auto-tuning mismatch                                   | Matmul, Conv2d, Depth-wise    | 1.66× (Matmul)           | TFLOPS vs. swept configs                                                      |
| RC2b: L2 cache residency                                     | Matmul $\geq 16384^2$         | —                        | L2 hit rate, DRAM throughput                                                  |
| RC3: TileLang LayerNorm spill                                | TileLang LayerNorm (large)    | —                        | <code>local_op</code> spill / occupancy (A100: 51.5 GB ld, 254 regs, 12% occ) |
| RC4: No Winograd                                             | Conv2d $3 \times 3$           | $\approx 2\text{--}3\%$  | SM utilization delta                                                          |

**Table 7: Evaluation heuristics on the A100-SXM4 (sm\_80, primary) and GH200 (sm\_90): optimized-kernel roofline fraction  $\rho$  (achieved/hardware peak; `mem`=HBM bandwidth, `comp`=FP16 Tensor-Core), PyTorch-relative efficiency  $E_{\text{lib}}$ , and the *cliff* (naive/optimized speedup). A high recovered  $\rho$  marks an *authoring artifact*; a low  $\rho$  that survives optimization marks a *structural residual*; the two architectures agree on every classification. Source: `experiments/results/NVIDIA_{A100-SXM4-40GB,GH200-480GB}/cliff_roofline.csv`.**

| Kernel                                                         | DSL      | Bound | A100-SXM4 (sm_80) |                   |                | GH200 (sm_90) |                  |                   |
|----------------------------------------------------------------|----------|-------|-------------------|-------------------|----------------|---------------|------------------|-------------------|
|                                                                |          |       | $\rho$            | $E_{\text{lib}}$  | Cliff          | $\rho$        | $E_{\text{lib}}$ | Cliff             |
| <i>Recovered family (authoring artifacts)</i>                  |          |       |                   |                   |                |               |                  |                   |
| LayerNorm                                                      | TileLang | mem   | 0.68              | 1.25              | 381×           | 0.59          | 1.20             | 1541×             |
| RMSNorm                                                        | TileLang | mem   | 0.70              | 10.1              | 324×           | 0.69          | 13.0             | 1520×             |
| MeanReduction                                                  | TileLang | mem   | 0.89              | 1.04              | 11.4×          | 0.87          | 0.97             | 13.9×             |
| BatchedMatmul                                                  | TileLang | mem   | 0.82              | 0.94              | 13.7×          | 0.86          | 1.11             | 20.2×             |
| MaxReduction                                                   | Triton   | mem   | 0.69              | 1.13              | 8.8×           | 0.59          | 1.20             | 6.4×              |
| MaxReduction                                                   | TileLang | mem   | 0.42              | 0.68              | 12.8×          | 0.41          | 0.84             | 16.2×             |
| LogSoftmax                                                     | TileLang | mem   | 0.55 <sup>§</sup> | 0.71 <sup>§</sup> | — <sup>§</sup> | 0.58          | 0.90             | 5.2×              |
| Softmax                                                        | TileLang | mem   | 0.55 <sup>§</sup> | 0.86 <sup>§</sup> | — <sup>§</sup> | 0.47          | 0.87             | 4.0× <sup>‡</sup> |
| <i>Structural residuals (survive best-effort optimization)</i> |          |       |                   |                   |                |               |                  |                   |
| Matmul                                                         | Triton   | comp  | 0.56              | 0.78              | 1.1×           | 0.47          | 0.68             | 1.3×              |
| Conv2d                                                         | TileLang | comp  | 0.34              | 0.59              | 5.6×           | 0.28          | 0.63             | 8.2×              |
| Conv2d                                                         | Triton   | comp  | 0.24              | 0.43              | 1.5×           | 0.16          | 0.34             | 2.5×              |
| Argmax                                                         | TileLang | mem   | 0.15              | 0.27              | 4.5×           | 0.10          | 0.22             | 3.8×              |

`torch.softmax` at this large shape, so its cliff is not a pure TileLang authoring comparison; the optimized  $\rho$  and  $E_{\text{lib}}$  are unaffected. <sup>§</sup> On the A100 stack (TileLang 0.1.6.post1) the in-tree TileLang Softmax/LogSoftmax variant’s 64 KB full-row shared-memory cache collapses occupancy ( $\rho=0.006/0.028$  in `cliff_roofline.csv`); we report the `opt_kernels` streaming variant, which reaches parity on both architectures ( $\rho=W/(tP)$  from its measured time,  $E_{\text{lib}}$  from `mitigation_retime.csv`); its cliff is omitted as the in-tree comparison does not transfer. `logsumexp` is omitted from the recovered family: its optimized kernel register-spills on sm\_90 (GH200;  $\rho < 0.02$ ,  $E_{\text{lib}} = 1.0\%$ ) but *not* on sm\_80 (A100: 0 spill, 93 regs,  $E_{\text{lib}} = 582\%$ , `ncu_summary.csv`)—an architecture-specific RC3-class residual discussed in section 7.5 and table 10.

re-timed kernel reaches 0.270 ms. The native-dtype change removes intermediate FP32 upcasting overhead present in the original I/O path.

This result does not represent a new optimization technique: `T.reduce` already existed in the API. The finding is that RC0 fully explains the normalization anomaly and is mechanically correctable without any algorithmic change.

**Table 8: Normalization kernel latency (ms, lower is better) before and after RC0 correction, re-timed on the A100-SXM4 (locked clocks).  $E_{\text{lib}}$  = library efficiency after fix, relative to PyTorch. *Before* is the unoptimized TileLang kernel (cf. table 3); *After* is the AKO4ALL-optimized kernel. The unoptimized kernels are memory-latency-bound (RC0) and hence clock-invariant, so *Before* matches the unlocked profile in table 3 to within noise. Input shape:  $8192 \times 8192$ .**

| Kernel           | PyTorch | Before | After        | $E_{\text{lib}}$  |
|------------------|---------|--------|--------------|-------------------|
| <i>TileLang</i>  |         |        |              |                   |
| LayerNorm (BF16) | 0.335   | 364    | <b>0.270</b> | <b>124%</b>       |
| RMSNorm (FP16)   | 2.52    | 294    | <b>0.254</b> | 990% <sup>†</sup> |

<sup>†</sup> RMSNorm  $E_{\text{lib}} > 100\%$  because the fixed kernel fuses the scaling pass, reducing total memory traffic relative to PyTorch’s unfused two-pass eager implementation; the fixed LayerNorm (124%) also slightly exceeds cuDNN’s already-fused `F.layer_norm`.

### 7.3 Convolution: Partial Recovery via Implicit GEMM

*Finding: Restructuring Conv2d as an implicit GEMM with FP16 Tensor Core acceleration, input padding to 16-byte alignment (addressing RC1), and an expanded autotune configuration space (addressing RC2) yields a 2.57× latency reduction ( $32.1\text{ ms} \rightarrow 12.5\text{ ms}$ ) in the AKO4ALL optimization campaign. On the A100-SXM4, where cuDNN is substantially faster, the same implicit-GEMM restructuring holds 36–76% of PyTorch/cuDNN efficiency across the  $1 \times 1\text{--}7 \times 7$  filter family (all correct)—the gap is narrowed but not closed.* The implicit GEMM restructuring unfolds the convolution input via on-the-fly tile formation (following the NHWC implicit GEMM approach [30]), converting the strided spatial access pattern into a dense matrix multiplication that Tensor Core units can accelerate in FP16. Input padding to a 16-byte boundary enables the compiler to emit LDG.128 loads on the channel dimension, partially recovering the vectorization deficit (RC1). The expanded autotune space adds smaller BLOCK\_K values (32, 16) for  $3 \times 3$  filters and additional pipeline stage counts, covering configurations absent from the default list (RC2).

Of the residual  $\approx 20\%$  gap, Winograd selection contributes only  $\approx 2\text{--}3\%$  (measured: a 0.06% deterministic-vs-non-deterministic delta on the A100-SXM4 for  $3 \times 3$  stride-1). The remaining  $\approx 17\text{--}18\%$  reflects general cuDNN implicit-GEMM tuning that the single-lowering Triton kernel does not match—kernel selection across data layouts, split-K, and shared-memory tile sizing (RC4). Closing this tuning gap, and adding Winograd support within a DSL kernel, are identified as future work.

**Table 9: Convolution mitigation across the filter family on the A100-SXM4 (optimized implicit-GEMM Triton, FP16; input  $32 \times 256 \times 128 \times 128$ ).  $E_{\text{lib}}$  = library efficiency relative to PyTorch/cuDNN; the mitigation holds 36–76% across the  $1 \times 1\text{--}7 \times 7$  filter family, all correct.**

| Filter       | Baseline $E_{\text{lib}}$ | Optimized $E_{\text{lib}}$ |                |
|--------------|---------------------------|----------------------------|----------------|
| $1 \times 1$ | 28.5%                     | <b>75.5%</b>               | Measured under |
| $3 \times 3$ | 19.3%                     | 42.4%                      |                |
| $5 \times 5$ | 13.7%                     | 38.5%                      |                |
| $7 \times 7$ | 9.8%                      | 36.1%                      |                |

locked clocks (graphics/memory 1215 MHz); locked-clock re-timing over 100 runs shows 0.0–0.9% run-to-run relative std-dev on near-parity kernels, superseding the prior cross-session clock-drift caveat (which reflected unlocked clocks).

### 7.4 GEMM and Reduction Kernels

Beyond the normalization kernels, we optimized the GEMM kernel and the full TileLang reduction/softmax family to test how broadly—and how completely—the RQ1 gaps recover.

**Square matmul (Triton, FP16).** Adding `@triton.autotune` with an expanded set of tile configurations and a `GROUP_SIZE_M` L2 cache swizzle (which reorders output tiles to improve L2 data reuse across the  $M$ -dimension) reduces latency from 1.08 ms to 0.79 ms ( $1.37\times$ ,  $E_{\text{lib}} = 77\%$ ) on a  $4096 \times 4096$  matmul, re-timed on the A100-SXM4; the same expanded search recovers  $E_{\text{lib}}$  from 32.5% to 81.8% ( $2.52\times$ ) at the RQ1  $16384^2$  shape. The original AKO4ALL campaign reported  $1.66\times / 108\%$  at  $4096^2$  on the RTX 4000 Ada, where the matmul mitigation slightly exceeded cuBLAS; on the faster A100 cuBLAS the gap narrows but does not close. Iterations 2–6 (persistent kernel, expanded configs, `max_num_imprecise_acc`) converged without further gain, confirming that the expanded configuration set and L2 swizzle are the effective ceiling for this shape. This confirms RC2 as a contributor to the GEMM gap and demonstrates that search-space expansion substantially narrows it for square shapes. This reconciles with the earlier default-grid finding that the stock 12-config autotune grid recovers  $\Delta \approx 0\text{pp}$ : the gain here is a property of the *expanded* search space (sweeping `GROUP_SIZE_M`, `num_warps`, and `num_stages`), not of a different problem shape.

**Argmax (TileLang, FP16).** Replacing global-memory element accesses with tiled shared-memory bulk loads and native FP16 I/O reduces latency from 11.97 ms to 2.31 ms ( $5.2\times$ ) on a  $8192 \times 32768$  reduction, re-timed on the A100-SXM4 (the original campaign reported  $9.26\times$  on the RTX 4000 Ada). The optimized kernel reaches  $E_{\text{lib}} = 27\%$ : it removes most of the unoptimized-TileLang overhead, but the A100’s native `torch.argmax` (0.62 ms) remains  $3.7\times$  faster, so parity is not reached on this GPU. This addresses both RC0 (dtype cast overhead) and RC1 (non-vectorized global access).

**Reduction and softmax family (TileLang).** The same RC0 fix—replacing `T.serial` reduction loops with native

**T.reduce** and streaming the reduction dimension in tiles (an online, max-rescaled pass for the softmax variants) rather than materializing the full row in a fragment—generalizes across the *entire* TileLang reduction and normalization family on the A100-SXM4 (table 10), not just the kernels of the original campaign. Every catastrophically slow kernel recovers to PyTorch-comparable or better: **max\_reduction**, **mean\_reduction**, and **logsumexp** *exceed* PyTorch (132%, 99%, 582%), and **softmax** and **log\_softmax** reach 86% and 71%—all numerically correct, with 4–29× speedups over the idiomatic kernels. This is the central evidence that the large normalization and reduction gaps in RQ1 are *authoring artifacts rather than DSL ceilings*: the idiomatic **T.serial** kernel is up to 300× slower, while on the A100 the one-pattern **T.reduce** rewrite matches or beats the vendor library on every kernel in the family. The lone within-A100 exception is index-returning **argmax** (27%), where the index bookkeeping is intrinsically heavier and **torch.argmax** is already well tuned—note that the *value-only* reduction over the identical shape (**max\_reduction**) reaches 132%, isolating the residual to the index pass. Re-timing the same optimized kernels on the GH200 confirms the pattern transfers—LayerNorm 120%, RMSNorm 1303%, MeanReduction 97%, MaxReduction 84%, Softmax 87%, LogSoftmax 90% (tables 7 and 10)—with one instructive exception: **logsumexp**, whose A100-tuned wide fp32 fragment register-spills on **sm\_90** (255 registers, ~280 GB of local-memory spill traffic, 12% occupancy by Nsight Compute, **ncu.summary.csv**), collapsing to 1.0%. The spill is an RC3-class effect, not the RC0 reduction idiom: retuning the fragment width recovers it only to 63%, so on the GH200 **logsumexp** is a genuine residual. That an “optimized” kernel validated on one GPU can be 100× off on another, with no correctness signal, is precisely the evaluation gap this paper argues a single-target benchmark cannot close.

## 7.5 Summary and Remaining Gap

Table 10 summarizes the per-category mitigation results, re-timed on the A100-SXM4, with optimized-kernel efficiency also reported for the GH200. The outcomes separate cleanly into two kinds, and the distinction is the central result of RQ3. The TileLang *normalization and reduction family* are *authoring artifacts*: correcting RC0 (**T.serial**→**T.reduce** plus native-dtype, tiled streaming I/O) brings every kernel to PyTorch-comparable or better—LayerNorm 124%, RMSNorm 990%, and the full reduction/softmax family at 71–582% (table 10)—closing gaps that were as large as 300× in RQ1. These are not DSL ceilings: once the kernel is written the competent way, the DSL matches or beats the vendor library on the whole family on the A100, and on the GH200 for every family kernel but **logsumexp**—whose optimized config register-spills on **sm\_90** (table 10), a reminder that a fixed kernel still carries a hardware-specific tuning that must be re-validated per target. The remaining gaps are *genuine residuals* that survive best-effort optimization: Conv2d (42%), large square GEMM (Matmul 77% vs cuBLAS), and index-returning Argmax (27%)—here even the fully-optimized DSL

**Table 10: Summary of mitigation results. A100-SXM4 latencies in ms (lower is better);  $E_{lib}$  = optimized-kernel efficiency relative to PyTorch/cuDNN on each GPU. Bold =  $\geq 95\%$  library efficiency on that GPU.**

| Kernel          | A100-SXM4 (ms) |        |              | $E_{lib}$         |                    | Root cause |
|-----------------|----------------|--------|--------------|-------------------|--------------------|------------|
|                 | PyTorch        | Before | After        | A100              | GH200              |            |
| <i>Triton</i>   |                |        |              |                   |                    |            |
| Matmul          | 0.607          | 1.081  | 0.788        | 77.0%             | 68%                | RC2        |
| Conv2d          | 3.44           | 17.86  | 8.12         | 42.4%             | 34%                | RC1+RC2    |
| <i>TileLang</i> |                |        |              |                   |                    |            |
| LayerNorm       | 0.335          | 364    | <b>0.270</b> | <b>124%</b>       | <b>120%</b>        | RC0        |
| RMSNorm         | 2.52           | 294    | <b>0.254</b> | 990% <sup>†</sup> | 1303% <sup>†</sup> | RC0        |
| Softmax         | 0.539          | 2.86   | 0.626        | 86.1%             | 87%                | RC0        |
| LogSoftmax      | 0.442          | 2.65   | 0.624        | 70.8%             | 90%                | RC0        |
| MaxReduction    | 0.560          | 11.97  | <b>0.418</b> | <b>131.6%</b>     | 84%                | RC0        |
| MeanReduction   | 0.768          | 10.65  | <b>0.766</b> | <b>99.5%</b>      | <b>97%</b>         | RC0        |
| LogSumExp       | 2.40           | 4.18   | <b>0.412</b> | <b>581.7%</b>     | 1.0% <sup>‡</sup>  | RC0/RC3    |
| Argmax          | 0.622          | 11.97  | 2.306        | 27.0%             | 22%                | RC0+RC1    |

<sup>†</sup> RMSNorm  $E_{lib} > 100\%$  due to kernel fusion; see table 8. <sup>‡</sup>

LogSumExp is the one family kernel whose optimized config does *not* transfer to the GH200: its A100-tuned wide fp32 fragment register-spills on **sm\_90** (255 regs, ~280 GB local-memory spill traffic, 12% occupancy; GH200 **ncu.summary.csv**) → 1.0%; even retuned to the GH200-optimal fragment width it caps at 63%, an RC3-class residual rather than an authoring artifact. On the A100 (**sm\_80**) the same kernel does *not* spill (0 local-memory traffic, 93 regs, 30.5% occupancy; A100 **ncu.summary.csv**), confirming the 581.7% is genuine and the spill is **sm\_90**-specific. A100 *ms* columns are re-timed on the A100-SXM4 (locked clocks); the memory-latency-bound norm/reduction *Before* kernels are clock-invariant and match the unlocked profile. GH200  $E_{lib}$  is the optimized-kernel efficiency from table 7 (unlocked). *Before* is the unoptimized DSL kernel (Matmul: plain Triton; Conv2d: baseline vs. implicit-GEMM Triton at  $3 \times 3$ , cf. table 9); Matmul at 4096<sup>2</sup>.

kernel does not reach the A100’s vendor library, because on the data-center A100 cuBLAS, cuDNN, and **torch.argmax** are themselves much faster than on the RTX 4000 Ada where the original campaign ran (where the matmul mitigation even reached 108%). Of the residual Conv2d gap, only  $\approx 2$ –3% is attributable to absent Winograd selection (RC4); the rest reflects general cuDNN implicit-GEMM tuning (kernel selection across data layouts, split-K, and shared-mem tile sizing) that the Triton kernel does not match.

The normalization, reduction, and convolution mitigations in table 10 are re-timed on the A100-SXM4 (locked clocks); the matmul figure is the AKO4ALL-campaign value pending the §7.3 reconciliation. The optimized kernels were revalidated against the same ViperBench tolerances as the unmitigated kernels (5/5 pass): the core mechanism (**T.serial**→**T.reduce** plus native-dtype I/O) reproduces correctly, and the edge-case correctness suite (NaN/Inf/denormal/all-equal) shows 0 crashes.

The normalization recoveries, while numerically dramatic, reflect correction of a known compiler deficiency (incorrect use of **T.serial** instead of **T.reduce**) rather than a novel optimization technique. They confirm that RC0 is both the



primary cause of TileLang’s normalization anomaly and mechanically correctable without algorithmic change. The convolution result (42%  $E_{\text{lib}}$  at  $3 \times 3$  after RC1+RC2, 36–76% across the filter family) establishes a tighter upper bound on what can be achieved without Winograd support on the A100, and leaves  $\approx 17$ –18pp of  $E_{\text{lib}}$  attributable to general cuDNN implicit-GEMM tuning that DSL kernels do not match (kernel selection across layouts, split-K, shared-mem tiling), and  $\approx 2$ –3pp to absent Winograd selection (RC4); both are candidates for future compiler work, with RC2b-style search-space expansion the higher-leverage priority.

## 8 Discussion

The empirical results suggest that current GPU DSLs are not uniformly competitive with vendor libraries; rather, their effectiveness depends strongly on the operator class and the optimizations available in the compiler stack. This has two implications. First, the observed gaps identify concrete weaknesses in present DSL systems, particularly for convolution. Second, the results provide practical guidance for users deciding when DSL kernels are likely to be a good substitute for library calls. We discuss these implications for DSL developers and practitioners, and then outline the main limitations of the study.

### 8.1 Implications for DSL Developers

Our results point to three practical directions for improving current GPU DSL systems.

**Improve vectorization for convolution accesses (requires compiler change).** RC1 (section 6) shows that convolution performance is limited in part by the failure to emit wide LDG.128 loads for strided spatial accesses. A direct compiler remedy is to extend alias and layout analysis to recognize the convolution access pattern and emit vectorized loads when the contiguous dimension is aligned (e.g., multiples of 8 for FP16). This change targets code generation rather than the programming model, and therefore does not require changes to the user-facing API.

**Make auto-tuning shape-aware (partly user-space fixable today; broader search requires compiler change).** RC2 (section 6) indicates that static tuning configurations transfer poorly from GEMM-like workloads to convolution. The RC2a portion is fixable in user space today: manually expanding the configuration list (e.g., `GROUP.SIZE.M/num_warps/num_stages`) recovers most of the GEMM gap, as shown in section 7. Closing the residual RC2b gap calls for auto-tuning strategies that are more adaptive to operator shape and memory-access structure. Prior work such as Ansor [29] shows that sketch-based search can discover effective schedules for irregular operators, including convolution, without relying on manually enumerated templates. A similar shape-aware search mechanism for Triton or TileLang would be a concrete step toward reducing this gap.

**Support algorithmic alternatives for convolution (requires compiler change).** RC4 highlights a limitation that cannot be addressed through local schedule tuning alone:

current DSL implementations do not expose algorithmic choices comparable to those available in cuDNN. In particular, support for Winograd-style transformations would allow the compiler or runtime to choose among direct convolution, GEMM lowering, and Winograd-based execution depending on filter shape and problem size. Although this requires a larger system change than RC1 or RC2, and the wall-clock benefit attributable specifically to Winograd is small ( $\approx 2$ –3%), it contributes to closing the remaining gap on common  $3 \times 3$  convolutions.

The results also have immediate consequences for users who adopt DSLs to replace library-backed kernels.

#### Performance depends strongly on operator class.

DSL kernels should not be assumed to match library performance by default. In our study, Triton is near parity on element-wise and normalization workloads, and can therefore offer a favorable trade-off between programmability and performance in these cases. For GEMM, the observed cost relative to cuBLAS may be acceptable when customization or fusion is required. Convolution is different: both Triton and TileLang remain substantially behind cuDNN in the evaluated settings. For convolution-heavy models, DSL kernels should therefore be treated as an optimization target that must be validated against a library baseline, not as a drop-in replacement.

**The root-cause analysis can guide debugging.** The mitigation results suggest a practical workflow for diagnosing underperforming DSL kernels. For TileLang reductions and normalization, developers should first replace serial `T.serial` accumulation with `T.reduce` (RC0a), since the dominant cost is exposed memory latency rather than thread synchronization; this is also where register pressure and spill should be inspected (RC3), which on our hardware appears in TileLang normalization rather than in convolution. For convolution, developers should check whether memory accesses are vectorized into wide LDG.128 loads (RC1), then evaluate whether the tuning configuration space is too narrow for the target shape, distinguishing a config that is merely absent from the default grid (RC2a, fixable by expanding the search) from one that needs a broader search strategy (RC2b). For  $3 \times 3$  filters, Winograd selection (RC4) is worth considering, but only as a small ( $\approx 2$ –3%) contributor. This sequence provides a compact checklist for performance debugging in Triton and TileLang.

**Benchmarks should gate on performance, not just correctness.** The passes-but-slow result (section 5.2) carries a concrete recommendation for benchmark designers: a correctness gate is necessary but not sufficient, and a kernel benchmark that reports speedup without gating on it certifies the wrong property. A lightweight, baseline-independent roofline check (section 7.1) is a practical acceptance criterion that needs no curated fast reference, and would have flagged every passes-but-slow kernel in our study.

## 8.2 Limitations

Our conclusions should be interpreted within the scope of the study.

**Operator scope.** We evaluate forward-pass kernels only. Backward kernels introduce different computation and memory patterns, including larger temporary state and more complex accumulation behavior, and may expose additional bottlenecks not captured here.

**Hardware scope.** Our primary measurements span two architectures, the NVIDIA A100-SXM4-40GB (Ampere, sm\_80) and the NVIDIA GH200 (Grace Hopper, sm\_90), with cross-architecture replay on the A100-PCIE-40GB (same sm\_80) and H100-80GB-HBM3 confirming that the gaps generalize across form factors and GPU families. The results nonetheless characterize NVIDIA datacenter platforms rather than all accelerators; in particular, we do not evaluate ROCm or Intel XPU backends, although these are important targets for future work.

**Software snapshot.** Both Triton and TileLang are evolving systems. The measurements in this paper reflect the software versions listed in section 3.6; later compiler or runtime changes may reduce some of the gaps we report. Our released benchmark suite is intended to support such re-evaluation.

## 9 Related Work

This section situates our work in the literature on GPU kernel DSLs (section 9.1), convolution optimization algorithms (section 9.2), performance analysis tools (section 9.3), and GPU benchmarking frameworks (section 9.4). Our study is the first to show that prevailing DSL and LLM-generated-kernel benchmarks admit performance-poor kernels, to characterize the resulting hidden gap with a hardware-grounded root-cause taxonomy, and to distill evaluation heuristics that flag poor kernels absent a comprehensive benchmark.

### 9.1 GPU Kernel DSLs

Halide [18] introduced the influential separation of algorithm specification from scheduling, enabling systematic search over loop transformations (tiling, vectorization, parallelism) without modifying the functional description. This principle informs both TVM [2] and TileLang [26].

Triton [23] departed from the explicit-schedule model in favor of an implicit tile abstraction in which the compiler infers shared memory layouts, synchronization, and pipelining. Its integration into PyTorch [1] has made it the most widely deployed GPU DSL. ThunderKittens [22] targets the opposite extreme: a thin C++ header library that exposes warp-level tile operations directly, achieving state-of-the-art attention throughput on H100 through explicit warp specialization. Our study is complementary: we characterize performance at the DSL level rather than at the hand-optimized C++ level.

TVM [2] and its auto-scheduler Ansor [29] demonstrate that hierarchical program search substantially outperforms template-guided auto-tuning for irregular operator shapes, including convolution. This is consistent with RC2 in our

analysis (section 6): the convolution search space is irregularly shaped relative to GEMM, and static configuration lists are insufficient. MLIR-based code generation [7, 24] offers a compiler-infrastructure approach to the same problem, generating near-peak-performance GEMM and fused-attention code through parametric Linalg transformations.

### 9.2 Convolution Optimization

Zhou et al. [30] provide a thorough characterization of the implicit GEMM convolution algorithm on GPU Tensor Cores, identifying the advantages of NHWC layout and on-the-fly tile formation over explicit im2col. Their analysis motivates the baseline construction in section 3.2.

Winograd convolution [9] reduces arithmetic by up to  $4\times$  for  $3\times 3$  filters at the cost of additional data transformation passes. In our measurements (section 7.3), the wall-clock benefit attributable specifically to Winograd selection is much smaller ( $\approx 2\text{--}3\%$ ); the arithmetic reduction does not translate proportionally because cuDNN’s general implicit-GEMM tuning dominates the gap. cuDNN’s selection of Winograd for appropriate shapes is a key component of its performance lead, as analyzed in section 6. Adding Winograd support as a first-class lowering pass within a DSL like Triton remains an open engineering problem; our analysis identifies it as future work (section 7.3).

### 9.3 Performance Analysis and Profiling

TritonForge [10] proposes a profiling-guided LLM loop for automated Triton kernel optimization, using Nsight Compute metrics to identify bottlenecks. Their finding that memory coalescing failures and low occupancy account for the majority of kernel-level inefficiencies is consistent with our RC1 and RC3 findings. The difference is scope: TritonForge focuses on automation of the fix-generation step, while our study focuses on systematic characterization across a kernel taxonomy.

Empirical performance studies of GPU libraries have examined cuDNN [4], cuBLAS [15], and CUTLASS [13] extensively, but performance comparisons *between* DSL and library implementations at the scale and granularity of our study are absent from the literature.

### 9.4 Benchmarking GPU Software

The benchmarks closest to our setting evaluate DSL and LLM-generated kernels. TritonBench [11] evaluates Triton kernel generation by LLMs, measuring functional correctness and GPU efficiency as a fraction of hardware peak—a roofline-anchored metric that predates and motivates the anchor we adopt in section 7.1, and which we build on rather than introduce. KernelBench [17] is the de-facto benchmark for LLM kernel generation, but it gates on correctness alone and reports speedup only as an outcome; its single reward-hacking guard flags only implausibly fast kernels, a weakness automated generators have exploited. Both evaluate only a narrow slice of the (DSL, data type, shape, hardware) space. Our work differs in object and claim: we do not propose another benchmark, but show that the prevailing ones admit

performance-poor kernels (section 5), characterize the resulting hidden gap with hardware counters, and distill evaluation heuristics and optimization patterns that guide development absent a comprehensive benchmark.

MLPerf Inference [19] benchmarks end-to-end inference throughput but treats the computation stack as a black box. Our study operates at the kernel level to attribute performance differences to specific compiler behaviors.

## 10 Threats to Validity

**Construct validity.** Our primary metric, library efficiency, measures DSL performance relative to the strongest library baseline we were able to exercise for each kernel. Because cuBLAS and cuDNN select among multiple internal algorithms, any failure to invoke their best-performing configuration would make the DSL appear closer to library performance than it actually is. To reduce this risk, we use library-backed PyTorch paths consistent with standard practice, invoke `cudaFind` for convolution, record the selected algorithms, and allow cuBLAS to use a large workspace. We report under default cuDNN flags (no explicit `cuda.benchmark` toggle, no NHWC conversion); convolutions run in the PyTorch-default NCHW layout, and algorithm-selection variance is bounded by `cudaFind` defaults.

**Internal validity.** Profiling-based analysis may be affected by measurement noise and tool overhead. We mitigate this threat by separating end-to-end timing from counter collection, using repeated runs for both, and verifying that key Nsight Compute measurements remain stable across repetitions. All experiments are executed in isolation on a fixed software and hardware setup to reduce confounding effects from concurrent workloads or environmental variation. Locked-clock re-timing (graphics 1215 MHz / memory 1215 MHz) over 100 runs shows run-to-run relative std-dev of 0.0–0.9% on near-parity kernels, so the reported gaps are statistically significant rather than artifacts of frequency variation.

**External validity.** Our study covers 22 kernels across five operator categories, but it does not represent the full design space of GPU workloads. In particular, the evaluated configurations emphasize common deep learning operators and shapes rather than uncommon cases such as highly dilated convolutions, extreme group counts, or very small batches. Conv2d is evaluated across  $1\times 1$ ,  $3\times 3$ ,  $5\times 5$ ,  $7\times 7$ , depthwise, and strided variants; the original submission tabulated only  $3\times 3$  stride-1. In addition, our primary measurements span two architectures—the A100-SXM4-40GB (Ampere, sm\_80) and the GH200 (Grace Hopper, sm\_90)—with further cross-architecture replay on the A100-PCIE-40GB (same sm\_80, a form-factor robustness check) and H100-80GB-HBM3 (a cross-family generalization check); results on other vendor backends may still differ.

**Conclusion validity.** Our root-cause claims are based on a combination of performance measurements, hardware-counter evidence, and targeted mitigations. Although the mitigation results strengthen the causal interpretation, they do

not rule out additional compiler or runtime factors that were not isolated in this study. The conclusions should therefore be read as evidence-supported explanations for the observed gaps, rather than as an exhaustive account of all possible causes.

**The suite is a probe, not a proposed benchmark.** Our 22-kernel suite is deliberately not offered as a comprehensive performance benchmark—a central claim of this work (section 5) is that such a benchmark is combinatorially infeasible to build and maintain. The suite instead functions as a diagnostic *probe*: it is large enough to expose the evaluation gap (correct kernels that existing benchmarks admit yet run far below the library), to root-cause that gap across operator classes, and to validate the proposed heuristics, but it is not intended to certify any kernel as production-ready. The contribution is the evaluation methodology—the passes-but-slow demonstration, the authoring/code-generation/library-maturity taxonomy, and the comparability-plus-roofline heuristics—rather than the suite’s size.

**Heuristic assumptions.** The roofline anchor depends on an essential-work model (bytes moved or FLOPs) and an assumed hardware peak for each kernel; mis-estimating either shifts the achieved fraction  $\rho$ , and the two checks disagree in a judgment band near  $\rho \approx 0.5$  (section 7.1). We therefore use  $\rho$  as a decision signal alongside the comparability screen rather than as a precise efficiency figure, and we record the peak assumptions with each measurement (table 7). The suite magnitude is measured on the A100 and the evaluation-gap and roofline results on the GH200—two primary architectures carrying complementary evidence—with per-architecture magnitudes reported where they differ.

## 11 Conclusion

We studied how to *evaluate* DSL GPU kernels: how a developer or an automated generator can tell whether a functionally correct Triton or TileLang kernel is actually performant. Across 22 kernels in five operator categories, we showed that the benchmarks practitioners rely on—KernelBench and TritonBench—gate on correctness and admit performance-poor kernels: a naive but idiomatic TileLang kernel passes the gate while running more than  $300\times$  slower than PyTorch, and the gate’s reward-hacking guard never fires because it only watches for kernels that are too *fast*. Because a benchmark comprehensive enough to close this gap is combinatorially infeasible, we instead characterized the hidden gap and distilled it into practical guidance.

The hidden gap is highly structured. Most of it is an *authoring* artifact: correcting a single TileLang idiom (`T.serial`→`T.reduce` with native-dtype I/O) recovers the normalization and reduction family to vendor-library parity or better—on both GPUs we test, with a single diagnosed exception (`logsumexp`, whose optimized kernel register-spills on the GH200)—closing gaps as large as  $1500\times$  (the naive-to-optimized authoring cliff on the GH200). The remainder is a *genuine residual*—convolution and large GEMM—where even a best-effort DSL kernel trails cuDNN/cuBLAS; absent Winograd selection



contributes only  $\approx 2\text{--}3\%$  of the convolution gap, the rest reflecting mature implicit-GEMM tuning. We separate causes by where the fix belongs (user-space authoring, compiler code generation, library maturity) and validate that lightweight heuristics—a comparability screen against the vendor baseline, anchored by a baseline-independent roofline fraction—reliably tell well-written kernels from poor ones, flagging the residuals (roofline  $\leq 0.28$ ) while clearing the recovered family (roofline 0.4–0.9).

Current GPU DSLs offer a strong programming model, but a correct kernel is not yet a fast one, and existing benchmarks do not close that gap. Until performance-aware benchmarks and compiler support mature, the evaluation heuristics and optimization patterns in this paper give developers—and the kernel-generating tools increasingly producing DSL code—a practical way to judge and improve kernel quality.

## 12 Data Availability Statement

All data and scripts used in this study are publicly available in an anonymous repository archived on Zenodo at: <https://doi.org/10.5281/zenodo.19337858>. The repository includes the benchmark datasets, evaluation scripts, and a README file describing the contents and usage instructions.

## References

- [1] Jason Ansel, Edward Yang, Horace He, Natalia Gimelshein, Animesh Jain, Michael Voznesensky, Bin Bao, Peter Bell, David Berard, Evgeni Burovski, Geeta Chauhan, Anjali Chourdia, Will Constable, Alban Desmaison, Zachary DeVito, Elias Ellison, Will Feng, Jiong Gong, Michael Gschwind, Brian Hirsh, Sherlock Huang, Kshiteej Kalambarkar, Laurent Kirsch, Michael Lazos, Mario Lezcano, Yanbo Liang, Jason Liang, Yinghai Lu, C. K. Luk, Bert Maher, Yunjie Pan, Christian Puhres, Matthias Reso, Mark Saroufim, Marcos Yukio Siraichi, Helen Suk, Shunting Zhang, Michael Suo, Phil Tillet, Xu Zhao, Eikan Wang, Keren Zhou, Richard Zou, Xiaodong Wang, Ajit Mathews, William Wen, Gregory Chanan, Peng Wu, and Soumith Chintala. 2024. PyTorch 2: Faster Machine Learning Through Dynamic Python Bytecode Transformation and Graph Compilation. In *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2* (La Jolla, CA, USA) (*ASPLOS '24*). Association for Computing Machinery, New York, NY, USA, 929–947. doi:10.1145/3620665.3640366
- [2] Tianqi Chen, Thierry Moreau, Ziheng Jiang, Lianmin Zheng, Eddie Yan, Haichen Shen, Meghan Cowan, Leyuan Wang, Yuwei Hu, Luis Ceze, Carlos Guestrin, and Arvind Krishnamurthy. 2018. TVM: An Automated End-to-End Optimizing Compiler for Deep Learning. In *13th USENIX Symposium on Operating Systems Design and Implementation (OSDI 18)*. USENIX Association, Carlsbad, CA, 578–594. <https://www.usenix.org/conference/osdi18/presentation/chen>
- [3] Sharan Chetlur, Cliff Woolley, Philippe Vandermersch, Jonathan Cohen, John Tran, Bryan Catanzaro, and Evan Shelhamer. 2014. cuDNN: Efficient Primitives for Deep Learning. arXiv:1410.0759 [cs.NE] <https://arxiv.org/abs/1410.0759>
- [4] Sharan Chetlur, Cliff Woolley, Philippe Vandermersch, Jonathan Cohen, John Tran, Bryan Catanzaro, and Evan Shelhamer. 2014. cuDNN: Efficient Primitives for Deep Learning. (2014). arXiv:1410.0759 [cs.NE] <https://arxiv.org/abs/1410.0759>
- [5] Tri Dao. 2023. FlashAttention-2: Faster Attention with Better Parallelism and Work Partitioning. arXiv:2307.08691 [cs.LG] <https://arxiv.org/abs/2307.08691>
- [6] Pieter Hijma, Stijn Heldens, Alessio Sclocco, Ben van Werkhoven, and Henri E. Bal. 2023. Optimization Techniques for GPU Programming. *ACM Comput. Surv.* 55, 11, Article 239 (March 2023), 81 pages. doi:10.1145/3570638
- [7] Navdeep Katel, Vivek Khandelwal, and Uday Bondhugula. 2022. MLIR-based code generation for GPU tensor cores. In *Proceedings of the 31st ACM SIGPLAN International Conference on Compiler Construction* (Seoul, South Korea) (*CC 2022*). Association for Computing Machinery, New York, NY, USA, 117–128. doi:10.1145/3497776.3517770
- [8] Chris Lattner, Mehdi Amini, Uday Bondhugula, Albert Cohen, Andy Davis, Jacques Pienaar, River Riddle, Tatiana Shpeisman, Nicolas Vasilache, and Oleksandr Zinenko. 2021. MLIR: Scaling Compiler Infrastructure for Domain Specific Computation. In *2021 IEEE/ACM International Symposium on Code Generation and Optimization (CGO)*. 2–14. doi:10.1109/CGO51591.2021.9370308
- [9] Andrew Lavin and Scott Gray. 2016. Fast Algorithms for Convolutional Neural Networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.
- [10] Haonan Li, Keyu Man, Partha Kanuparth, Hanning Chen, Wei Sun, Sreen Tallam, Chenguang Zhu, Kevin Zhu, and Zhiyun Qian. 2025. TritonForge: Profiling-Guided Framework for Automated Triton Kernel Optimization. arXiv:2512.09196 [cs.SE] <https://arxiv.org/abs/2512.09196>
- [11] Jianling Li, Shangzhan Li, Zhenye Gao, Qi Shi, Yuxuan Li, Zefan Wang, Jiacheng Huang, Haojie Wang, Jianrong Wang, Xu Han, Zhiyuan Liu, and Maosong Sun. 2025. TritonBench: Benchmarking Large Language Model Capabilities for Generating Triton Operators. In *Findings of the Association for Computational Linguistics: ACL 2025*, Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (Eds.). Association for Computational Linguistics, Vienna, Austria, 23053–23066. doi:10.18653/v1/2025.findings-acl.1183
- [12] Sparsh Mittal and Jeffrey S. Vetter. 2014. A Survey of Methods for Analyzing and Improving GPU Energy Efficiency. *ACM Comput. Surv.* 47, 2, Article 19 (Aug. 2014), 23 pages. doi:10.1145/2636342

- [13] NVIDIA Corporation. 2017. *CUTLASS: CUDA Templates and Python DSLs for High-Performance Linear Algebra*.
- [14] NVIDIA Corporation. 2024. NVIDIA Nsight Compute. <https://developer.nvidia.com/nsight-compute>.
- [15] NVIDIA Corporation. 2026. *cuBLAS: Basic Linear Algebra on NVIDIA GPUs*.
- [16] OpenAI. 2021. Introducing Triton: Open-Source GPU Programming for Neural Networks. <https://openai.com/index/triton/>.
- [17] Anne Ouyang, Simon Guo, Simran Arora, Alex L. Zhang, William Hu, Christopher Ré, and Azalia Mirhoseini. 2025. KernelBench: Can LLMs Write Efficient GPU Kernels? arXiv:2502.10517 [cs.LG] <https://arxiv.org/abs/2502.10517>
- [18] Jonathan Ragan-Kelley, Connelly Barnes, Andrew Adams, Sylvain Paris, Frédo Durand, and Saman Amarasinghe. 2013. Halide: a language and compiler for optimizing parallelism, locality, and recomputation in image processing pipelines. *SIGPLAN Not.* 48, 6, 519–530. doi:10.1145/2499370.2462176
- [19] Vijay Janapa Reddi, Christine Cheng, David Kanter, Peter Mattson, Guenther Schmuelling, Carole-Jean Wu, Brian Anderson, Maximilien Breughe, Mark Charlebois, William Chou, Ramesh Chukka, Cody Coleman, Sam Davis, Pan Deng, Greg Diamos, Jared Duke, Dave Fick, J. Scott Gardner, Itay Hubara, Sachin Idgunji, Thomas B. Jablin, Jeff Jiao, Tom St. John, Pankaj Kanwar, David Lee, Jeffery Liao, Anton Lokhmotov, Francisco Massa, Peng Meng, Paulius Micikevicius, Colin Osborne, Gennady Pekhimenko, Arun Tejusve Raghunath Rajan, Dilip Sequeira, Ashish Sirasao, Fei Sun, Hanlin Tang, Michael Thomson, Frank Wei, Ephrem Wu, Lingjie Xu, Koichi Yamada, Bing Yu, George Yuan, Aaron Zhong, Peizhao Zhang, and Yuchen Zhou. 2020. MLPerf Inference Benchmark. In *2020 ACM/IEEE 47th Annual International Symposium on Computer Architecture (ISCA)*. 446–459. doi:10.1109/ISCA45697.2020.00045
- [20] sebastienwood. 2022. Example Conv2D in Triton. GitHub Discussion #591, <https://github.com/triton-lang/triton/discussions/591>. Accessed: 2026-03-25.
- [21] Md. Maruf Hossain Shuvo, Syed Kamrul Islam, Jianlin Cheng, and Bashir I. Morshed. 2023. Efficient Acceleration of Deep Learning Inference on Resource-Constrained Edge Devices: A Review. *Proc. IEEE* 111, 1 (2023), 42–91. doi:10.1109/JPROC.2022.3226481
- [22] Benjamin Frederick Spector, Simran Arora, Aaryan Singhal, Arjun Parthasarathy, Daniel Y Fu, and Christopher Re. 2025. ThunderKittens: Simple, Fast, and  $\text{\textit{Adorable}}$  Kernels. In *The Thirteenth International Conference on Learning Representations*. <https://openreview.net/forum?id=0fJfVOSUra>
- [23] Philippe Tillet, H. T. Kung, and David Cox. 2019. Triton: an intermediate language and compiler for tiled neural network computations. In *Proceedings of the 3rd ACM SIGPLAN International Workshop on Machine Learning and Programming Languages (Phoenix, AZ, USA) (MAPL 2019)*. Association for Computing Machinery, New York, NY, USA, 10–19. doi:10.1145/3315508.3329973
- [24] Nicolas Vasilache, Oleksandr Zinenko, Aart J. C. Bik, Mahesh Ravishanker, Thomas Raoux, Alexander Belyaev, Matthias Springer, Tobias Gysi, Diego Caballero, Stephan Herhut, Stella Laurenzo, and Albert Cohen. 2022. Composable and Modular Code Generation in MLIR: A Structured and Retargetable Approach to Tensor Compiler Construction. arXiv:2202.03293 [cs.PL] <https://arxiv.org/abs/2202.03293>
- [25] Guibin Wang, YiSong Lin, and Wei Yi. 2010. Kernel Fusion: An Effective Method for Better Power Efficiency on Multithreaded GPU. In *2010 IEEE/ACM Int'l Conference on Green Computing and Communications & Int'l Conference on Cyber, Physical and Social Computing*. 344–350. doi:10.1109/GreenCom-CPSCom.2010.102
- [26] Lei Wang, Yu Cheng, Yining Shi, Zhiwen Mo, Zhengju Tang, Wenhao Xie, Tong Wu, Lingxiao Ma, Yuqing Xia, Jilong Xue, Fan Yang, and Zhi Yang. 2026. TileLang: Bridge Programmability and Performance in Modern Neural Kernels. In *The Fourteenth International Conference on Learning Representations*. <https://openreview.net/forum?id=Jb1WkNSfUB>
- [27] Samuel Williams, Andrew Waterman, and David Patterson. 2009. Roofline: an insightful visual performance model for multicore architectures. *Commun. ACM* 52, 4 (April 2009), 65–76. doi:10.1145/1498765.1498785
- [28] M. Wolfe. 1989. More iteration space tiling. In *Proceedings of the 1989 ACM/IEEE Conference on Supercomputing (Reno, Nevada, USA) (Supercomputing '89)*. Association for Computing Machinery, New York, NY, USA, 655–664. doi:10.1145/76263.76337
- [29] Lianmin Zheng, Chengfan Jia, Minmin Sun, Zhao Wu, Cody Hao Yu, Ameer Haj-Ali, Yida Wang, Jun Yang, Danyang Zhuo, Koushik Sen, Joseph E. Gonzalez, and Ion Stoica. 2020. Ansor: Generating High-Performance Tensor Programs for Deep Learning. In *14th USENIX Symposium on Operating Systems Design and Implementation (OSDI 20)*. USENIX Association, 863–879. <https://www.usenix.org/conference/osdi20/presentation/zheng>
- [30] Yangjie Zhou, Mengtian Yang, Cong Guo, Jingwen Leng, Yun Liang, Quan Chen, Minyi Guo, and Yuhao Zhu. 2021. Characterizing and Demystifying the Implicit Convolution Algorithm on Commercial Matrix-Multiplication Accelerators. In *2021 IEEE International Symposium on Workload Characterization (IISWC)*. 214–225. doi:10.1109/IISWC53511.2021.00029